

CSE1500 – WEB AND DATABASE TECHNOLOGY

DB LECTURE 4

MORE ON SQL

Alessandro Bozzon

cse1500-ewi@tudelft.nl

AT THE END OF THIS LECTURE, YOU SHOULD BE ABLE TO....

- ▶ **Describe** and **design** SQL programs for the retrieval of data from tables
- ▶ **Describe** and **design** SQL programs for the creation, altering, and manipulation of tables
- Develop** logical database schema, with principled design that enforce data integrity
- Prototype** and **deploy** database applications using open-source database systems (e.g., PostgreSQL)

EXAMPLE DATABASES

EXAMPLE DB1: EMPLOYEES

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

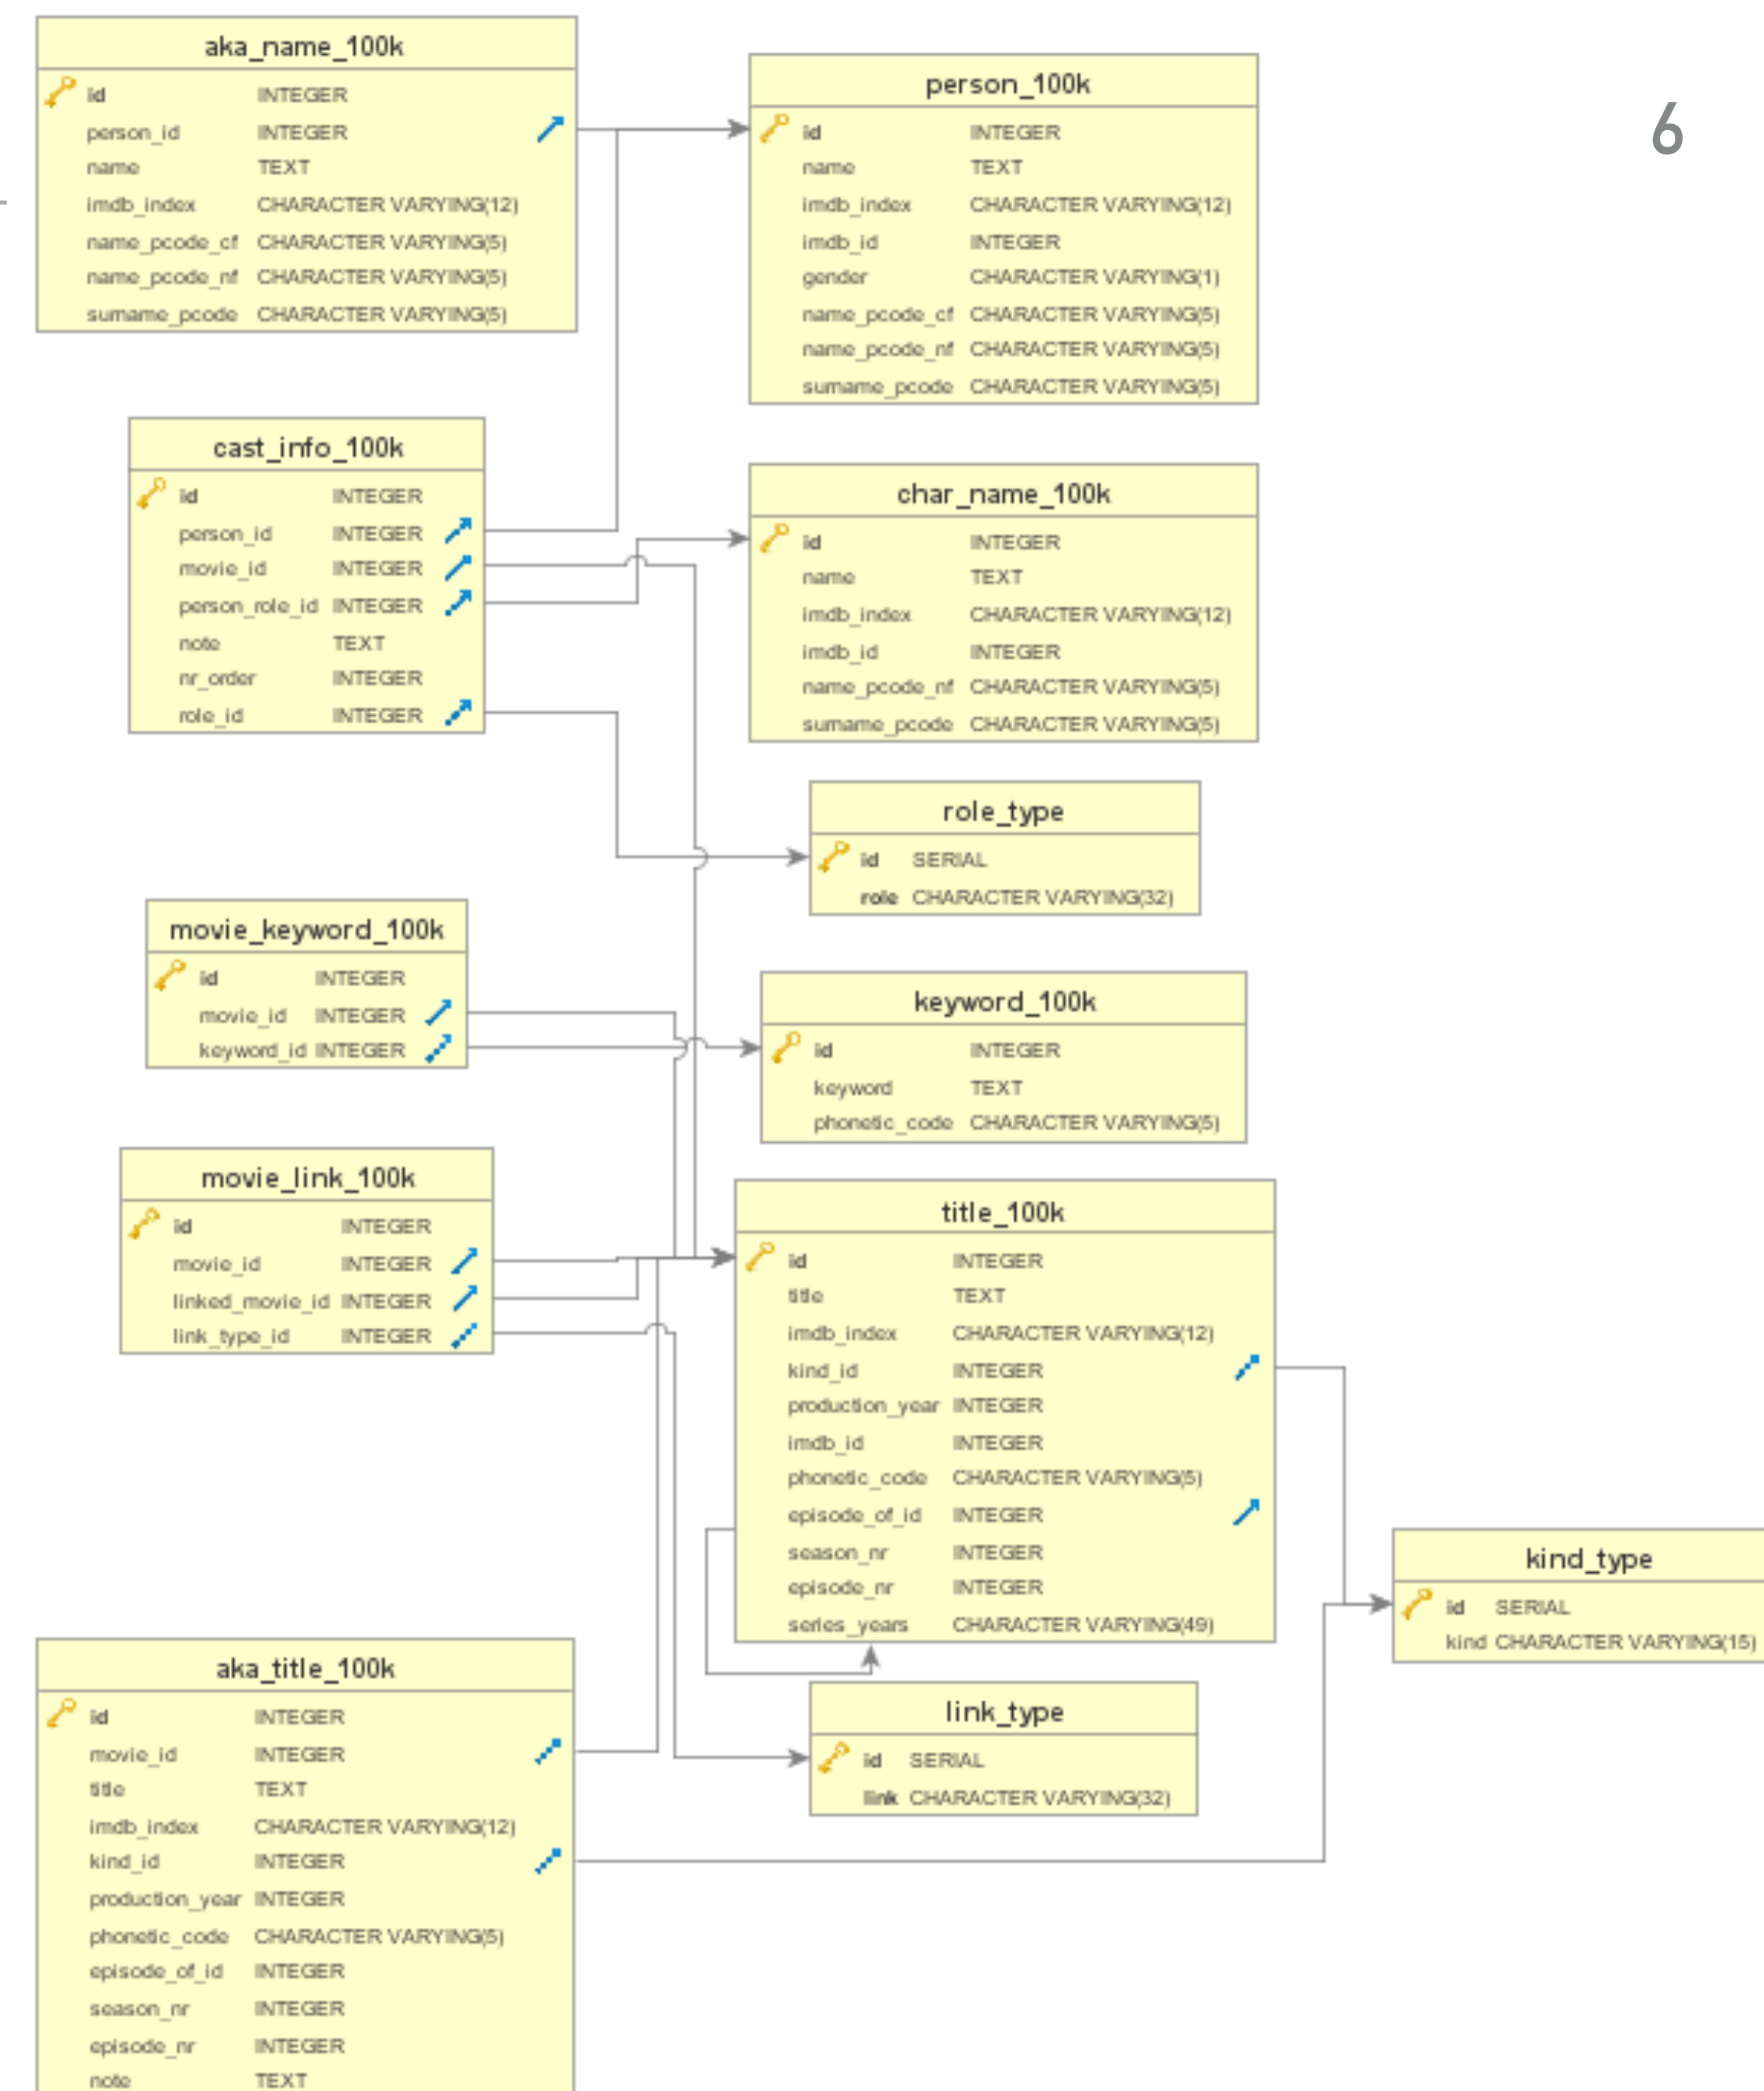
Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

EXAMPLE DB2: PRODUCTS

Supplier				Supply			Products				
<u>CodeS</u>	NameS	Shareholders	Office	<u>CodeS</u>	<u>CodeP</u>	Amount	<u>CodeP</u>	NameP	Color	Size	Storehouse
S1	John	2	Amsterdam	S1	P1	300	P1	Sweater	Red	40	Amsterdam
S2	Victor	1	Den Haag	S1	P2	200	P2	Jeans	Green	48	Den Haag
S3	Anna	3	Den Haag	S1	P3	400	P3	Shirt	Blu	48	Rotterdam
S4	Angela	2	Amsterdam	S1	P4	200	P4	Shirt	Blu	44	Amsterdam
S5	Paul	3	Utrecht	S1	P5	100	P5	Skirt	Blu	40	Den Haag
				S1	P6	100	P6	Coat	Red	42	Amsterdam
				S2	P1	300					
				S2	P2	400					
				S3	P2	200					
				S4	P3	200					
				S4	P4	300					
				S4	P5	400					

EXAMPLE DB3: IMDB

- ▶ A subset of the schema and data from the [IMDB.com](https://www.imdb.com) website
 - ▶ Actors (person_100k), Movies (title_100k), and Actors in Movies (cast_info_100k)
 - ▶ Plus aliases, keywords, movie genres, etc.
- ▶ We will use MongoDB and Neo4J implementations of the same database (obviously, with different schemas)
- ▶ Get it (with import instructions) here
 - ▶ https://docs.google.com/document/d/1jj3cMAnk6Rc0mHkkOAIYDzYLjKisCuyj4-3KF9l-_8o



AGGREGATE QUERIES

AGGREGATE QUERIES

- ▶ **Aggregate Query:** query in which the result depends on the consideration of **sets of rows**
- ▶ The result is a single (**aggregated**) value
- ▶ Expressed in the SELECT clause
 - ▶ aggregate operators are evaluated on the rows accepted by the WHERE conditions
- ▶ SQL92 offers five aggregate operators
 - ▶ COUNT, SUM, MAX, MIN, AVG
- ▶ Except for COUNT, these functions return a NULL value when no rows are selected

OPERATOR COUNT

- ▶ COUNT returns the number of rows or distinct values

```
COUNT (<* | [DISTINCT | ALL] TargetList >)
```

- ▶ The DISTINCT keyword forces the count of distinct values in the attribute list


COUNT EXAMPLE /1

► Find the **number of suppliers** in the database

```
SELECT COUNT (*)  
FROM Supplier
```



Supplier			
CodeS	NameS	Shareholders	Office
S1	John	2	Amsterdam
S2	Victor	1	Den Haag
S3	Anna	3	Den Haag
S4	Angela	2	Amsterdam
S5	Paul	3	Utrecht



count
5


COUNT EXAMPLE /2

► Find the **number of suppliers** with at least one supply

```
SELECT COUNT (*)  
FROM Supply
```

- Is it right?
- Equivalent to `SELECT COUNT(CodeP)`
or `SELECT COUNT(CodeS)`

Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400

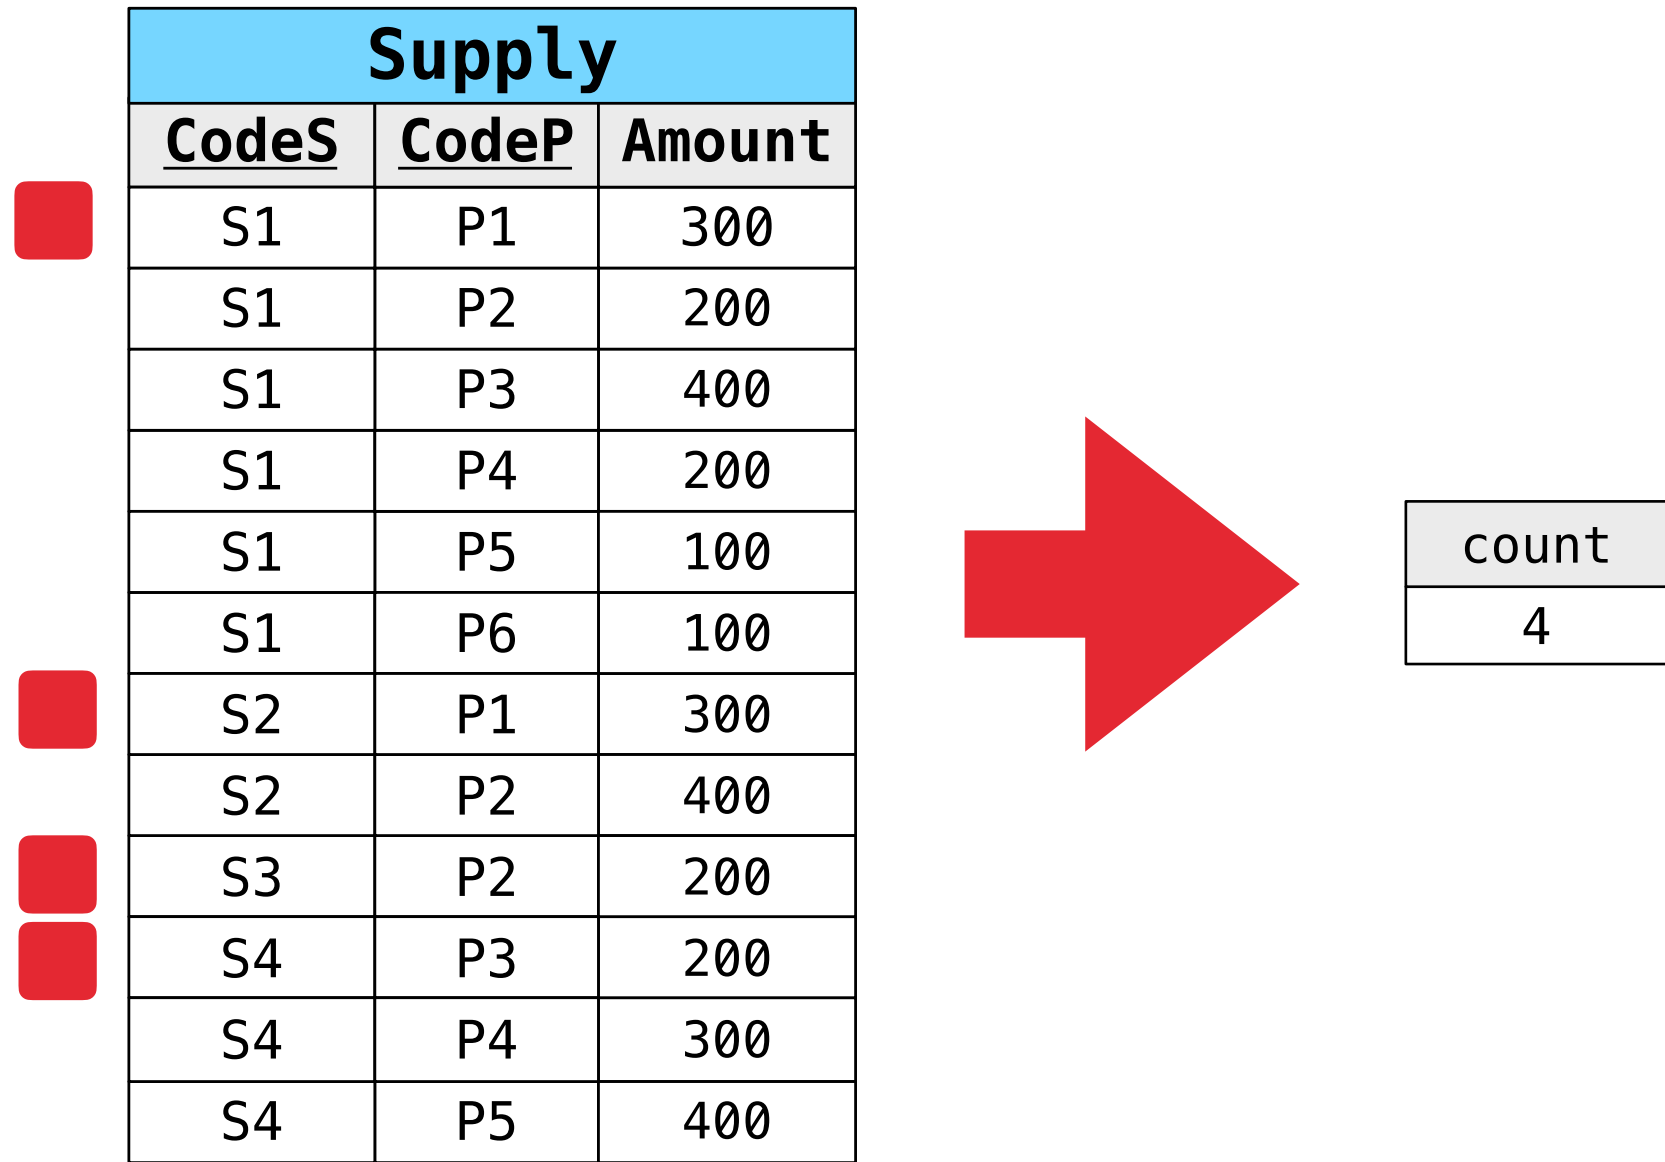


count
12

COUNT EXAMPLE /3

► Find the **number of suppliers** with at least one supply

```
SELECT COUNT (DISTINCT CodeS)
FROM   Supply
```



COUNT EXAMPLE /4

► **Count** the number of suppliers that supply the product “P2”

```
SELECT COUNT(*)  
FROM Supply  
WHERE CodeP = 'P2'
```

► Is it right?

Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400

count

3

OPERATORS SUM,MAX,MIN,AVG


- ▶ SUM,MAX,MIN,AVG
 - ▶ Allowed arguments are attributes or expressions
- ▶ SUM,AVG
 - ▶ Only numeric types
- ▶ MAX,MIN
 - ▶ Attribute must be sortable
 - ▶ Applied also on strings and timestamps

SUM EXAMPLE


► Find the **total** number of supplied items for product “P2”

```
SELECT SUM(Amount)
FROM Supply
WHERE CodeP = 'P2'
```

Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400



CodeS	CodeP	Amount
S1	P2	200
S2	P2	400
S3	P2	200



count
800

NULL VALUES AND AGGREGATES

- ▶ All aggregate operations ignore tuples with NULL values on the aggregated attributes
 - ▶ COUNT: number of input rows for which the value of expression is not NULL
 - ▶ SUM, AVG, MAX, MIN: NULL values are not considered
- ▶ The COALESCE function can be used to force a value for NULL

```
SELECT AVG(season_nr)
FROM title_100k
```

```
SELECT AVG(COALESCE(season_nr, 1))
FROM title_100k
```


AGGREGATE QUERY AND TARGET LIST

- ▶ This is an incorrect query, although syntactically admissible

```
SELECT FirstName, Surname, MAX(Salary)
FROM Employee JOIN Department ON Dept = DeptName
WHERE Department.City = 'London'
```


- ▶ Whose name? The target list must be homogeneous
- ▶ The GROUP BY clause will help us

GROUPING ROWS


- ▶ Queries may apply aggregate operators to subsets of rows
- ▶ **For each product** find the total amount of supplied items

```
SELECT CodeP, SUM(Amount)
FROM Supply
GROUP BY CodeP
```

Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400



CodeS	CodeP	Amount
S1	P1	300
S2	P1	300
S1	P2	200
S2	P2	400
S3	P2	200
S1	P3	400
S4	P3	200
S1	P4	200
S4	P4	300
S1	P5	100
S4	P5	400
S1	P6	100



CodeP	Amount
P1	600
P2	800
P3	600
P4	500
P5	500
P6	100

GROUP BY CLAUSE /1

- ▶ The order of the grouping attributes does not matter
- ▶ The SELECT clause can contain
 - ▶ Attributes specified in the GROUP BY clause
 - ▶ Aggregated functions
 - ▶ Attributes univocally determined by attributes already specified in the GROUP BY clause

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

```
SELECT Office
FROM Employee
GROUP BY Dept
```

▶ Incorrect Query

GROUP BY CLAUSE /2

- ▶ The order of the grouping attributes does not matter
- ▶ The SELECT clause can contain
 - ▶ Attributes specified in the GROUP BY clause
 - ▶ Aggregated functions
 - ▶ Attributes univocally determined by attributes already specified in the GROUP BY clause

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

```
SELECT DeptName,D.City,COUNT(*)
FROM Employee E JOIN Department D ON E.Dept=D.DeptName
GROUP BY DeptName
```

▶ Incorrect Query

GROUP BY CLAUSE /3

- ▶ The order of the grouping attributes does not matter
- ▶ The SELECT clause can contain
 - ▶ Attributes specified in the GROUP BY clause
 - ▶ Aggregated functions
 - ▶ Attributes univocally determined by attributes already specified in the GROUP BY clause

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

```
SELECT DeptName,D.City,COUNT(*)
FROM Employee E JOIN Department D ON E.Dept=D.DeptName
GROUP BY DeptName, D.City
```


▶ Correct Query

GROUPING ROWS


- ▶ Queries may apply aggregate operators to subsets of rows
- ▶ **For each product** sold by suppliers in Den Haag, find the total amount of supplied items

```
SELECT CodeP, SUM(Amount)
FROM Supply JOIN Supplier ON Supply.CodeS = Supplier.CodeS
WHERE Office = 'Den Haag'
GROUP BY CodeP
```

Supplier				Supply		
CodeS	NameS	Shareholders	Office	CodeS	CodeP	Amount
S1	John	2	Amsterdam	S1	P1	300
S1	John	2	Amsterdam	S1	P2	200
S1	John	2	Amsterdam	S1	P3	400
S1	John	2	Amsterdam	S1	P4	200
S1	John	2	Amsterdam	S1	P5	100
S1	John	2	Amsterdam	S1	P6	100
S2	Victor	1	Den Haag	S2	P1	300
S2	Victor	1	Den Haag	S2	P1	300
S2	Victor	1	Den Haag	S2	P2	400
S3	Anna	3	Den Haag	S3	P2	200
S4	Angela	2	Amsterdam	S4	P3	200
S4	Angela	2	Amsterdam	S4	P4	300
S4	Angela	2	Amsterdam	S4	P5	400



CodeP	Amount
P1	300
P1	300
P2	200
P2	400



CodeP	Amount
P1	600
P2	600

HAVING CLAUSE /1

- ▶ Conditions on the result of an aggregate operator require the HAVING clause
- ▶ Only predicates containing aggregate operators should appear in the argument of the HAVING clause
- ▶ Find the departments in which the average salary of employees working in office number 20 **is higher than 25**


```
SELECT Dept
FROM Employee
WHERE Office = '20'
GROUP BY Dept
HAVING AVG(Salary) > 25
```


HAVING CLAUSE /2


- Find the total number of supplied items for products that count **at least 600 total supplied items**

```
SELECT CodeP, SUM(Amount)
FROM Supply
GROUP BY CodeP
HAVING SUM(Amount) >= 600
```

Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400



CodeS	CodeP	Amount
S1	P1	300
S2	P1	300
S1	P2	200
S2	P2	400
S3	P2	200
S1	P3	400
S4	P3	200
S1	P4	200
S4	P4	300
S1	P5	100
S4	P5	400
S1	P6	100




CodeP	Amount
P1	600
P2	800
P3	600

HAVING CLAUSE /3

- Find the code of red products supplied by more than one supplier

```
SELECT Supply.CodeP
FROM Supply JOIN Products ON Supply.CodeP = Product.CodeP
WHERE Color = 'Red'
GROUP BY Supply.CodeP
HAVING COUNT(*) > 1
```

Products					Supply		
CodeP	NameP	Color	Size	Storehouse	CodeS	CodeP	Amount
P1	Sweater	Red	40	Amsterdam	S1	P1	300
P1	Sweater	Red	40	Amsterdam	S2	P1	300
P2	Jeans	Green	48	Den Haag	S1	P2	200
P2	Jeans	Green	48	Den Haag	S2	P2	400
P2	Jeans	Green	48	Den Haag	S3	P2	200
P3	Shirt	Blu	48	Rotterdam	S1	P3	400
P3	Shirt	Blu	48	Rotterdam	S4	P3	200
P4	Shirt	Blu	44	Amsterdam	S1	P4	200
P4	Shirt	Blu	44	Amsterdam	S4	P4	300
P5	Skirt	Blu	40	Den Haag	S1	P5	100
P5	Skirt	Blu	40	Den Haag	S4	P5	400
P6	Coat	Red	42	Amsterdam	S1	P6	100



CodeP
P1

NESTED QUERIES

SUBQUERIES

- ▶ A parenthesised SELECT-FROM-WHERE statement can be used as a value in a number of places, including FROM, WHERE, and HAVING clauses
 - ▶ *subquery* or *nested* query
 - ▶ Example: in place of a table in the FROM clause, we can have another query, and then assert a condition on its results
- ▶ A common use of subqueries is to perform tests for **set membership**, **set comparisons**, and **set cardinality**
- ▶ The use of nested queries may produce less declarative queries, but they often improve readability
 - ▶ Complex queries can become very difficult to understand

THE IN OPERATOR

- ▶ `<tuple> IN <relation>` is true if and only if the tuple is a member of the relations
 - ▶ `<tuple> NOT IN <relation>` means the opposite
- ▶ IN expression can appear in WHERE clauses
 - ▶ The relation is often a subquery
 - ▶ e.g. `WHERE firstName IN (SELECT lastName FROM Employee)`
 - ▶ It is possible to specify a pre-defined list of values
 - ▶ e.g. `WHERE firstName IN ('Jan', 'Arie', 'Henk')`
- ▶ It allows **problem decomposition**, typically with a “bottom-up” approach

EXAMPLE /1

- ▶ Find the *name* of **all** the suppliers **of** product "P2"

TWO SUBPROBLEMS

- ▶ CODE OF P2 SUPPLIERS
- ▶ NAME OF SUCH SUPPLIERS

WITH JOIN

```
SELECT NameS
FROM   Supplier, Supply
WHERE  Supplier.CodeS = Supply.CodeS AND CodeP = 'P2'
```

WITH IN AND NESTED QUERIES

```
SELECT NameS
FROM   Supplier
WHERE  CodeS IN (SELECT CodeS
                  FROM   Supply
                  WHERE  CodeP = 'P2')
```

EXAMPLE /2

- Find the *name* of supplier of **at least one red product**

THREE SUBPROBLEMS

- ▶ **CODE OF RED PRODUCTS**
- ▶ **CODE OF SUPPLIERS THAT SUPPLY SUCH PRODUCTS**
- ▶ **NAME OF SUCH SUPPLIERS**

WITH JOIN

```
SELECT NameS
FROM Supplier, Supply, Products
WHERE Supplier.CodeS = Supply.CodeS AND Supply.CodeP = Products.CodeP
AND Color = 'Red'
```

WITH IN AND NESTED QUERIES

```
SELECT NameS
FROM Supplier
WHERE CodeS IN (SELECT CodeS
FROM Supply
WHERE CodeP IN (SELECT CodeP
FROM Products
WHERE Color = 'Red'))
```

NESTED QUERIES THAT RETURN ONE TUPLE

- ▶ If a subquery is guaranteed to produce one tuple, then the result of the subquery can be used as a **value**
 - ▶ Typically, a single tuple is guaranteed by key constraints of attributes SELECTed by the subquery
- ▶ A run-time error occurs if there is no tuple or more than one tuple
- ▶ Usually, the tuple has one attribute, but with a tuple constructor we might have many => **row subquery**

EXAMPLE /1

- Find the *code* of suppliers having office in the same city as S1

WITH JOIN

```
SELECT Su2.CodeS
FROM   Supplier AS Su1, Supplier AS Su2
WHERE  Su1.Office = Su2.Office AND Su1.CodeS = 'S1'
```

WITH NESTED QUERIES

```
SELECT CodeS
FROM   Supplier
WHERE  Office = (SELECT Office
                  FROM   Supplier
                  WHERE  CodeS = 'S1')
```


EXAMPLE /2

- ▶ Find the code of suppliers with less shareholders than the supplier having the **maximum** number of shareholders

WITH NESTED
QUERIES

```
SELECT CodeS
FROM Supplier
WHERE Shareholders < (SELECT MAX(Shareholders)
                      FROM Supplier)
```

WITH JOIN??

MAYBE. THINK ABOUT IT :)

EXAMPLE /3

- Find the name of the suppliers that supply at least one product supplied by suppliers of red products

Difficult to express with joins

- ▶ Code of red products
- ▶ Code of suppliers that supply such products
- ▶ Code of products supplied by suppliers of red products
- ▶ Code of the suppliers of the products supplied by suppliers of red products
- ▶ Name of such suppliers

```
SELECT NameS
FROM Supplier
WHERE CodeS IN (SELECT CodeS
FROM Supply
WHERE CodeP IN (SELECT CodeP
FROM Supply
WHERE CodeS IN (SELECT CodeS
FROM Supply
WHERE CodeP IN (SELECT CodeP
FROM Product
WHERE Color = 'Red')))))
```

EXAMPLE /3 WITH JOIN

- Find the name of the suppliers that supply at least one product supplied by suppliers of red products

```
SELECT NameS
FROM Supplier, Supply AS SA, Supply AS SB, Supply AS SC, Products
WHERE Supplier.CodeS = SA.CodeS AND
      SA.CodeP = SB.CodeP AND
      SB.CodeS = SC.CodeS AND
      SC.CodeP = Products.CodeP AND
      Products.Color = 'Red'
```

EXAMPLE WITH NOT IN

- ▶ Find the name of the suppliers that **DO NOT** supply product P2
- ▶ Can you express it with JOIN?

```
SELECT NameS
FROM   Supplier, Supply
WHERE  Supplier.CodeS = Supply.CodeS AND CodeP <> 'P2'
```

EXAMPLE WITH NOT IN

- ▶ Find the name of the suppliers that **DO NOT** supply product P2
- ▶ Can you express it with JOIN?

```
SELECT NameS
FROM Supplier, Supply
WHERE Supplier.CodeS = Supply.CodeS AND CodeP <> 'P2'
```

- ▶ **WRONG!!!**
- ▶ The SQL above would answer the query:
- ▶ Find the name of the suppliers that supply **at least one product different from** P2

EXAMPLE WITH NOT IN/2

- ▶ Find the name of the suppliers that **DO NOT** supply product P2
- ▶ We need to exclude from the result set suppliers that supply **P2**

```
SELECT NameS
FROM Supplier
WHERE CodeS NOT IN (SELECT CodeS
                     FROM Supply
                     WHERE CodeP = 'P2')
```

EXAMPLE WITH NOT IN /3

- ▶ Find the name of the suppliers that supply **ONLY** product P2
- ▶ We need to exclude from the result set suppliers that supply products different from P2
- ▶ But only keep the ones that actually supply something

```
SELECT NameS
FROM Supplier
WHERE CodeS NOT IN (SELECT CodeS
                    FROM Supply
                    WHERE CodeP <> 'P2')
AND CodeS IN (SELECT CodeS
              FROM Supply)
```


ON QUANTIFIERS

▶ Existential Quantifiers: **Easy!** :)

▶ Find the name of the suppliers that supply **at least one product different from** P2

▶ Universal Quantifiers: **Hard!** :(

▶ Find the name of the suppliers that **DO NOT** supply product P2

▶ Find the name of the suppliers that supply **ONLY** product P2

▶ Find the name of the suppliers that **DO NOT** supply red products

SCOPE OF ATTRIBUTES, VARIABLES, AND CORRELATED QUERIES

- ▶ A subquery can use attributes and/or variables defined by outermost queries in their WHERE clause
 - ▶ this is sometimes referred to as "*transfer of bindings*"
- ▶ The two queries are said to be correlated
- ▶ Semantics: the nested query is evaluated **for each row of the external query**

THE EXISTS OPERATOR/1

- Find the name of the suppliers that supplied P2 at least once

```
SELECT NameS
FROM Supplier
WHERE EXISTS (SELECT *
               FROM Supply
               WHERE CodeP='P2' AND Supplier.CodeS = Supply.CodeS)
```

- !!! We need to test the existence of a supply for the evaluation of suppliers
- Suppliers.CodeS = Supply.CodeS imposes a correlation between external and internal query

THE EXISTS OPERATOR /2

- ▶ Find all the homonyms
- ▶ i.e., people with the same first name and surname, but different BSN

```
SELECT *  
FROM Person P1  
WHERE EXISTS (  
    SELECT *  
    FROM Person P2  
    WHERE P1.FirstName = P2.FirstName AND  
          P1.Surname = P2.Surname AND  
          P1.BSN <> P2.BSN)
```

LIMITATIONS

- ▶ A query cannot refer to attributes in a subquery, or in a query at the same level of nesting

```
SELECT *  
FROM Employee  
WHERE Dept IN (SELECT DeptName  
                FROM Department D1  
                WHERE DeptName = 'Production') OR  
             Dept IN (SELECT DeptName  
                      FROM Department D2  
                      WHERE D2.City = D1.City)
```

- ▶ The query is incorrect because variable D1 is not visible in the second nested query

EXAMPLE OF NOT EXISTS OPERATOR /1

- ▶ Find the name of the suppliers that **DO NOT** supply P2 or
- ▶ Find the name of the suppliers for whom it **does not exist at least one** supply of P2

```
SELECT NameS
FROM Supplier
WHERE NOT EXISTS (SELECT *
                  FROM Supply
                  WHERE CodeP = 'P2' AND
                        Supplier.CodeS = Supply.CodeS)
```

EXAMPLE OF NOT EXISTS OPERATOR /2

- Find all the persons who **do not have** homonyms

```
SELECT *  
FROM Person P1  
WHERE NOT EXISTS (SELECT *  
                  FROM Person P2  
                  WHERE P1.FirstName = P2.FirstName AND  
                        P1.Surname = P2.Surname AND  
                        P1.BSN <> P2.BSN)
```

TUPLE CONSTRUCTOR

- ▶ The comparison with the nested query may involve more than one attributes
- ▶ The attributes must be enclosed within a pair of curved brackets (tuple constructor)
- ▶ The query in the previous slide can be expressed as:

```
SELECT *
FROM Person P1
WHERE (FirstName, Surname) NOT IN (SELECT FirstName, Surname
FROM Person P2
WHERE P1.BSN <> P2.BSN)
```


ANY OPERATOR /1

- ▶ Find the employees who work in departments in London

```
SELECT FirstName, Surname
FROM Employee
WHERE Dept = ANY (SELECT DeptName
                  FROM Department
                  WHERE City = 'London')
```

- ▶ But also

```
SELECT FirstName, Surname
FROM Employee, Department
WHERE Dept = DeptName AND Department.City = 'London'
```

ANY OPERATOR /2

- Find the employees of the Planning department, having the same first name as a member of the Production department

```
SELECT FirstName, Surname
FROM Employee
WHERE Dept = 'Planning' AND FirstName = ANY (SELECT FirstName
                                              FROM Employee
                                              WHERE Dept = 'Production')
```

- But also

```
SELECT E1.FirstName, E1.Surname
FROM Employee E1, Employee E2
WHERE E1.FirstName = E2.FirstName AND E2.Dept = 'Production' AND E1.Dept = 'Planning'
```

EXAMPLE MIN AND MAX

- ▶ Queries using the aggregate operators max and min can be expressed with nested queries

```
SELECT Dept
FROM Employee
WHERE Salary >= ALL (SELECT Salary
                     FROM Employee)
```

- ▶ But also

```
SELECT Dept
FROM Employee
WHERE Salary IN (SELECT MAX(Salary)
                FROM Employee)
```

EXAMPLE OF NEGATION WITH ALL

- Find the departments in which there is **no one** named Brown

```
SELECT DeptName
FROM Department
WHERE DeptName <> ALL (SELECT Dept
                        FROM Employee
                        WHERE Surname='Brown')
```

- Alternatively (more next)

```
SELECT DeptName
FROM Department
EXCEPT
SELECT Dept
FROM Employee
WHERE Surname='Brown'
```

NESTED AND AGGREGATED QUERIES /1

- ▶ Find the code of the suppliers that supply ALL products

```
SELECT CodeS
FROM Supply
GROUP BY CodeS
HAVING COUNT(*) = (SELECT COUNT(*)
                   FROM Products)
```

- ▶ All the products that can be supplied are in the Product table
- ▶ A supplier supply all the products if the number of distinct supplied product is the same as the number of available products

NESTED AND AGGREGATED QUERIES /2

- Find the code of the suppliers that supplied **as much as** S2

```
SELECT CodeS
FROM Supply
WHERE CodeP IN (SELECT CodeP
                FROM Supply
                WHERE CodeS = 'S2')

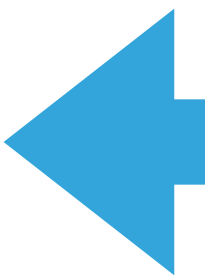
GROUP BY CodeS
HAVING COUNT(*) = (SELECT COUNT(*)
                  FROM Supply
                  WHERE CodeS = 'S2')
```

- The number of supplies of a supplier should be equal to the number of supplies by S2

SET QUERIES

SET QUERIES

- ▶ Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:
 - ▶ (subquery) **INTERSECT** [ALL] (subquery)
 - ▶ (subquery) **EXCEPT** [ALL] (subquery)
- ▶ (subquery) **UNION** [ALL] (subquery)



CAN BE EXPRESSED WITH
OTHER OPERATORS
(TYPICALLY SUB-QUERIES)



ENHANCEMENT OF THE
EXPRESSIVE POWER

SET SEMANTIC OF SET QUERIES

- ▶ Although the SELECT-FROM-WHERE statement uses **bag semantics**, the *default* for union, intersection, and difference is **set semantics**.
 - ▶ That is, duplicates are eliminated as the operation is applied
- ▶ Motivation: Efficiency
 - ▶ When projecting attributes, it is easier to avoid **eliminating duplicates**. Just work tuple-at-a-time.
 - ▶ When doing intersection or difference, it is most efficient to **sort the relations first**
At that point you may as well eliminate the duplicates anyway

UNION

- ▶ A single SELECT cannot represent unions of values from two or more tables

A SQL query snippet is displayed within a light blue rectangular frame. The text inside the frame is "A UNION [ALL] B". The word "UNION" is in blue, "[ALL]" is in red, and "A" and "B" are in black.

```
A UNION [ALL] B
```

- ▶ It executes the union of two relational expressions
 - ▶ Expressions generated by SELECT clauses
 - ▶ Table A and Table B must be **union compatible**
 - ▶ i.e. have compatible schema
 - ▶ same number of output fields, in the same order, and with the same or compatible data types
- ▶ Duplicate removal
 - ▶ UNION removes duplicates
 - ▶ UNION ALL does not remove duplicates

UNION EXAMPLE /1

- Find the code of red products **OR** products supplied by S2 (or both)

```
SELECT CodeP
FROM Products
WHERE Color = 'Red'
```

UNION


```
SELECT CodeP
FROM Supply
WHERE CodeS = 'S2'
```

UNION EXAMPLE /1

► Find the code of red products **OR** products supplied by S2 (or both)

Products				
CodeP	NameP	Color	Size	Storehouse
P1	Sweater	Red	40	Amsterdam
P2	Jeans	Green	48	Den Haag
P3	Shirt	Blu	48	Rotterdam
P4	Shirt	Blu	44	Amsterdam
P5	Skirt	Blu	40	Den Haag
P6	Coat	Red	42	Amsterdam


Supply		
CodeS	CodeP	Amount
S1	P1	300
S1	P2	200
S1	P3	400
S1	P4	200
S1	P5	100
S1	P6	100
S2	P1	300
S2	P2	400
S3	P2	200
S4	P3	200
S4	P4	300
S4	P5	400



CodeP
P1
P6



CodeP
P1
P2
P6



CodeP
P1
P2



UNION EXAMPLE /2

- ▶ Find the first names and surnames of the employees

```
SELECT FirstName AS Name  
FROM Employee
```

```
UNION
```

```
SELECT Surname  
FROM Employee
```

INTERSECTION

A diagram showing the SQL syntax for the intersection operator. The text 'A INTERSECT [ALL] B' is centered within a white rectangular box. This box is surrounded by a light blue border, which is itself inside a larger, slightly offset light blue rectangular frame.

```
A INTERSECT [ALL] B
```

- ▶ Intersection of two subqueries
 - ▶ returns all rows that are both in the result of A and in the result of B
- ▶ As for the UNION operator, schema must be **union compatible**
- ▶ Duplicate rows are eliminated unless INTERSECT ALL is used.
- ▶ Not supported by all RDBMS (e.g. not supported by MySQL)

INTERSECT EXAMPLE

- Find cities hosting both suppliers' offices and products' storehouses

```
SELECT Office  
FROM Supplier
```

INTERSECT

```
SELECT Storehouse  
FROM Product
```

- Equivalent to

```
SELECT DISTINCT Office  
FROM Supplier, Product  
WHERE Office = Storehouse
```

INTERSECT EXAMPLE /2

- Find the surnames of employees that are also first names

```
SELECT FirstName  
FROM Employee  
  
INTERSECT  
  
SELECT Surname  
FROM Employee
```

- Equivalent to

```
SELECT DISTINCT E1.FirstName  
FROM Employee E1, Employee E2  
WHERE E1.FirstName = E2.Surname
```


EXCEPT

```
A EXCEPT [ALL] B
```

- ▶ Difference set operator
 - ▶ Returns all rows that are in the result of A but not in the result of B
- ▶ As for the UNION operator, schema must be **union compatible**
- ▶ Not supported by all RDBMS (e.g. not supported by MySQL)

EXCEPT EXAMPLE

- Find cities hosting suppliers' offices, but not products' storehouses

```
SELECT Office
FROM Supplier

EXCEPT

SELECT Storehouse
FROM Products
```

- Can be represented with a nested query using the NOT IN operator

```
SELECT DISTINCT Office
FROM Supplier
WHERE Office NOT IN (SELECT Storehouse
                     FROM Products)
```

EXCEPT EXAMPLE /2

- Find the surnames of employees that are not also first names

```
SELECT FirstName  
FROM Employee
```

```
EXCEPT
```

```
SELECT Surname  
FROM Employee
```

SCHEMA DEFINITION

DEFINING A DATABASE SCHEMA

- ▶ A database schema comprises:
 - ▶ declarations for the relations ("tables") of the database
 - ▶ domains associated with each attribute
 - ▶ integrity constraints
- ▶ A schema has a *name* and an *owner* (the authorisation)
- ▶ Many other kinds of elements may also appear in the database schema, including:
 - ▶ privileges, views, indexes, triggers

- ▶ Syntax:

```
CREATE SCHEMA [ SchemaName ]  
[ [ authorisation ] Authorisation ]  
{ SchemaElementDefinition }
```

DOMAINS

- ▶ Specify the content of attributes
- ▶ Two categories
 - ▶ **Elementary** (predefined by the standard)
 - ▶ **User-defined** (not available in all RDBMs implementations)

```
CREATE DOMAIN Grade AS SMALLINT  
    DEFAULT NULL  
    CHECK (Grade >= 0 AND Grade <=10)
```

ELEMENTARY DOMAINS (DATA TYPES)/1

- ▶ Bit
 - ▶ Single boolean values or strings of boolean values (may be variable in length)
 - ▶ Syntax: BIT [varying] [(Length)]
- ▶ Exact numeric domains
 - ▶ Exact values, integer or with a fractional part
 - ▶ Four Alternatives
 - ▶ NUMERIC [(Precision [, Scale])]: fixed point number, with user-specified Precision digits, of which Scale digits to the right of decimal point.
 - ▶ DECIMAL [(Precision [, Scale])]: functionally equivalent to NUMERIC
 - ▶ INTEGER: a finite subset of the integers that is machine-dependent
 - ▶ SMALLINT: a machine-dependent subset of the integer domain type

ELEMENTARY DOMAINS (DATA TYPES)/2

▶ **Approximate real values**

- ▶ Based on floating point representation
- ▶ `FL0AT [(Precision)]`: floating point number, with user-specified precision of at least `n` digits. By default `n` is 53, but it can be less
- ▶ `REAL`: floating point numbers, with machine-dependent precision
- ▶ `Double Precision`: double-precision floating point numbers, with machine-dependent precision

ELEMENTARY DOMAINS (DATA TYPES)/3

▶ Temporal **Instants**

- ▶ DATE: format yyyy-mm-dd
- ▶ TIME [(Precision)] [with time zone]: format hh:mm:ss:p with an optional decimal point and fractions of a second following.
- ▶ TIMESTAMP [(Precision)] [with time zone]: format yyyy-mm-dd hh:mm:ss:p

▶ Temporal **intervals**

- ▶ INTERVAL FirstUnitOfTime [TO LastUnitOfTime]
- ▶ Units of time are divided into two groups:
 - ▶ year, month
 - ▶ day, hour, minute, second
- ▶ In PostgreSQL the syntax is different: interval '2 months ago'

ELEMENTARY DOMAINS (DATA TYPES)/4

- ▶ Geometric Types: two-dimensional spatial object
 - ▶ point, line, lseg, box, path, Open path, polygon, circle
- ▶ Network Address Types: to store IPv4, IPv6, and MAC addresses
 - ▶ cidr, inet, macaddr, macaddr8
- ▶ JSON Types
 - ▶ json: data is stored an exact copy of the input text
 - ▶ jsonb: data is stored in a decomposed binary format
- ▶ XML Type, used to store XML data
- ▶ Composite Types: represents the structure of a row or record
- ▶ UUID, Array, Ranges, Text Search (to support full text search)

TABLE DEFINITION

- ▶ An SQL table consists of
 - ▶ an ordered set of attributes
 - ▶ a (possibly empty) set of constraints
- ▶ Statement CREATE TABLE
 - ▶ defines a relation schema, creating an empty instance
- ▶ Constraints: integrity checks on attributes
- ▶ OtherConstraints: integrity constraints on the table
- ▶ Syntax:

```
CREATE TABLE TableName (  
    AttributeName Domain [DefaultValue] [Constraints]  
    {, AttributeName Domain [DefaultValue] [Constraints]}  
    [OtherConstraints]  
)
```

EXAMPLE OF CREATE TABLE

```
CREATE TABLE Employee (  
  RegNo      CHARACTER(6) PRIMARY KEY,  
  FirstName  CHARACTER(20) NOT NULL,  
  Surname    CHARACTER(20) NOT NULL,  
  Dept       CHARACTER(15)  
             REFERENCES Department(DeptName)  
             ON DELETE SET NULL  
             ON UPDATE CASCADE  
  Salary     DECIMAL (9) DEFAULT 0,  
  City       CHARACTER(15),  
  UNIQUE(Surname, FirstName)  
)
```

DEFAULT DOMAIN VALUES

- ▶ Define the value that the attribute must assume when a value is not specified during row insertion
- ▶ Syntax:

```
DEFAULT < GenericValue | USER | CURRENT_USER | SESSION_USER | SYSTEM_USER | NULL >
```

- ▶ `GenericValue` represents a value compatible with the domain, in the form of a constant or an expression
- ▶ `USER*` is the login name of the user who issues the command

CONSTRAINTS /1

- ▶ Constraints are conditions that must be verified by every database instance
- ▶ Defined in the CREATE or ALTER TABLE operations
- ▶ Automatically verified by the DB after each operation
- ▶ Advantages
 - ▶ declarative specification of constraints
 - ▶ unique centralised verification
- ▶ Disadvantages
 - ▶ might slow down execution
 - ▶ pre-defined type of constraints
 - ▶ e.g. no constraint on aggregated data
 - ▶ but triggers can help

CONSTRAINTS /2

- ▶ An operation that violates a constraints might cause two type of reactions:
 - ▶ the operation is **aborted**, causing an application error
 - ▶ a **compensation action** is taken, to reach a new consistent state
- ▶ Three type of constraints
 - ▶ Intra-relational constraints (or **table constraints**)
 - ▶ Inter-relational constraints (or **referential integrity constraints**)
 - ▶ *Generic* integrity constraints and **assertions**

INTRA-RELATIONAL CONSTRAINTS

Intra-relational constraints involve a single relation

- ▶ NOT NULL (on single attributes)
 - ▶ upon tuple insertion, the attribute MUST be specified
- ▶ UNIQUE: permits the definition of candidate keys
 - ▶ for single attributes: UNIQUE, after the domain
 - ▶ for multiple attributes: UNIQUE(Attribute , Attribute)
- ▶ PRIMARY KEY: defines the primary key
 - ▶ once for each table
 - ▶ implies NOT NULL
 - ▶ syntax like UNIQUE

PRIMARY KEY VS. UNIQUE

- ▶ The SQL standard allows DBMS implementers to make their own distinctions between PRIMARY KEY and UNIQUE
 - ▶ Example: some DBMS might automatically create an index (data structure to speed search) in response to PRIMARY KEY, but not UNIQUE
- ▶ However, standard SQL requires these distinctions:
 - ▶ There can be **only one** PRIMARY KEY for a relation, but several UNIQUE attributes
 - ▶ No attribute of a PRIMARY KEY can ever be NULL in any tuple.
 - ▶ But attributes declared UNIQUE may have NULLs
 - ▶ and there may be several tuples with NULL!

EXAMPLE OF INTRA-RELATIONAL CONSTRAINTS

- ▶ Each pair of FirstName and Surname uniquely identifies each element

```
FirstName CHARACTER(20) NOT NULL  
Surname CHARACTER(20) NOT NULL  
UNIQUE(FirstName, Surname)
```

- ▶ Note the difference with the following (stricter) definition

```
FirstName CHARACTER(20) NOT NULL UNIQUE  
Surname CHARACTER(20) NOT NULL UNIQUE
```

INTER-RELATIONAL CONSTRAINTS

- ▶ Constraints may take into account several relations
- ▶ REFERENCES and FOREIGN KEY permit the definition of referential integrity constraints. Syntax:
 - ▶ for single attributes: REFERENCES, after the domain
 - ▶ for multiple attributes: FOREIGN KEY (Attribute1 , Attribute2) REFERENCES Table (Attribute1 , Attribute2)
- ▶ **It is possible to associate reaction policies to violations of referential integrity**

REACTION POLICIES FOR REFERENTIAL INTEGRITY CONSTRAINTS

- ▶ Reactions operate **on** the referencing table, after changes **to** the referenced table
- ▶ Violations may be introduced
 - ▶ by updates on the referred attribute
 - ▶ by row deletions
- ▶ Reactions (can be specific to an event)
 - ▶ CASCADE: propagate the change
 - ▶ SET NULL: nullify the referring attribute
 - ▶ SET DEFAULT: assign the default value to the referring attribute
 - ▶ NO ACTION: forbid the change on the external table
- ▶ Syntax:

```
ON < DELETE | UPDATE > < CASCADE | SET NULL | SET DEFAULT | NO ACTION >
```

EXAMPLE OF INTER-RELATIONAL CONSTRAINT

```
CREATE TABLE Employee (  
  RegNo      CHARACTER(6) PRIMARY KEY,  
  FirstName  CHARACTER(20) NOT NULL,  
  Surname    CHARACTER(20) NOT NULL,  
  Dept       CHARACTER(15)  
             REFERENCES Department (DeptName)  
             ON DELETE SET NULL  
             ON UPDATE CASCADE  
  Salary     DECIMAL (9) DEFAULT 0,  
  City       CHARACTER(15),  
  UNIQUE (Surname, FirstName)  
)
```

Referencing

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

Referenced

Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

SCHEMA UPDATES

- ▶ Two SQL statements:
 - ▶ ALTER: to modify a domain, the schema of a table, or a user
 - ▶ DROP: to remove schema, domain, table, etc.

```
ALTER TABLE Department ADD COLUMN NoOfOffices NUMERIC(4)
```

```
ALTER TABLE Department ADD CONSTRAINT UniqueAddress UNIQUE(Address)
```

```
DROP TABLE TempTable CASCADE
```

RELATIONAL CATALOGUES

- ▶ The catalog contains:
 - ▶ The data dictionary
 - ▶ The description of the data contained in the data base (tables, etc.)
 - ▶ Statistics about the data (distribution, access, growth)
- ▶ It is based on a relational structure (reflexive)
- ▶ It can be queried!
- ▶ The SQL92 standard describes a Definition Schema (composed of tables) and an Information Schema (composed of views)

```
SELECT table_name
FROM information_schema.tables
WHERE table_schema = 'public'
```

DATA MANIPULATION

DATA MODIFICATION IN SQL

- ▶ Statements for:
 - ▶ insertion INSERT
 - ▶ deletion DELETE
 - ▶ change of attribute values UPDATE
- ▶ All the statements **can operate on a set of tuples** (set-oriented)
- ▶ In the condition it is possible to access other relations

INSERTIONS /1

```
INSERT INTO TableName [(AttributeList)] <VALUES(ListofValues)|SELECT SQL>
```

► Using Values

```
INSERT INTO Department(DeptName,City) VALUES ('Production','Toulouse')
```

► Using a subquery

```
INSERT INTO LondonProducts(  
    SELECT Code, Description  
    FROM Product  
    WHERE ProdArea = `London`  
)
```

INSERTIONS /2

- ▶ The ordering of the attributes (if present) and of values is meaningful (first value with the first attribute, and so on)
- ▶ If *AttributeList* is omitted, all the relation attributes are considered, in the order in which they appear in the table definition
- ▶ If *AttributeList* does not contain all the relation attributes, to the remaining attributes it is assigned:
 - ▶ the DEFAULT value (if defined)
 - ▶ the NULL value
 - ▶ PRIMARY KEYS might get special handling

DELETIONS /1

- ▶ The DELETE statement removes from the table all the tuples that satisfy the condition

```
DELETE FROM TableName [WHERE Condition]
```

- ▶ The removal may produce deletions from other tables if a referential integrity constraint with CASCADE policy has been defined
- ▶ If WHERE clause is omitted, DELETE removes all the tuples

DELETIONS /2

- ▶ To remove all the tuples from DEPARTMENT keeping the table

```
DELETE FROM Department
```

- ▶ To remove table DEPARTMENT completely (content and schema)

```
DROP TABLE Department CASCADE
```

- ▶ Remove the `Production` department

```
DELETE FROM Department WHERE DeptName = `Production`
```

- ▶ Remove the departments without employees

```
DELETE FROM Department WHERE DeptName NOT IN (SELECT Dept FROM Employee)
```

UPDATES /1

```
UPDATE TableName  
  SET Attribute = <Expression | SELECT SQL | NULL | default>  
  {, Attribute = <Expression | SELECT SQL | NULL | default>}  
  [WHERE Condition]
```

► Examples

```
UPDATE Employee  
  SET Salary = Salary + 5  
  WHERE FirstName = 'Mary' AND LastName = 'Brown'
```

```
UPDATE Employee  
  SET Salary = Salary * 1.1  
  WHERE Dept = 'Administration'
```

UPDATES /2

- ▶ Since the language is set oriented, the order of the statements is important

```
UPDATE Employee  
  SET Salary = Salary * 1.1  
 WHERE Salary <= 30
```

```
UPDATE Employee  
  SET Salary = Salary * 1.15  
 WHERE Salary > 30
```

- ▶ If the statements are issued in this order, some employees may get a double raise

WRAPPING UP

TODAY WE COVERED

- ▶ SQL as
 - ▶ a Retrieval Language
 - ▶ a Schema Creation and Modification Language
 - ▶ a Data Manipulation Language

END OF LECTURE