

CSE1500 – WEB AND DATABASE TECHNOLOGY

DB LECTURE 7

EVEN MORE ON SQL

Alessandro Bozzon

cse1500-ewi@tudelft.nl

AT THE END OF THIS LECTURE, YOU SHOULD BE ABLE TO....

- ▶ **Describe** and **design** SQL queries that make use of the WITH clause
- ▶ **Describe** and **design** SQL queries that make use of Views
- ▶ **Describe** and **design** constraints at application level with Triggers
- ▶ **Administer** users and permissions

EXAMPLE DATABASES

EXAMPLE DB1: EMPLOYEES

| Employee | | | | | |
|-----------|----------|----------------|--------|--------|----------|
| FirstName | Surname | Dept | Office | Salary | City |
| Mary | Brown | Administration | 10 | 45 | London |
| Charles | White | Production | 20 | 36 | Toulouse |
| Gus | Green | Administration | 20 | 40 | Oxford |
| Jackson | Neri | Distribution | 16 | 45 | Dover |
| Charles | Brown | Planning | 14 | 80 | London |
| Laurence | Chen | Planning | 7 | 73 | Worthing |
| Pauline | Bradshaw | Administration | 75 | 40 | Brighton |
| Alice | Jackson | Production | 20 | 46 | Toulouse |

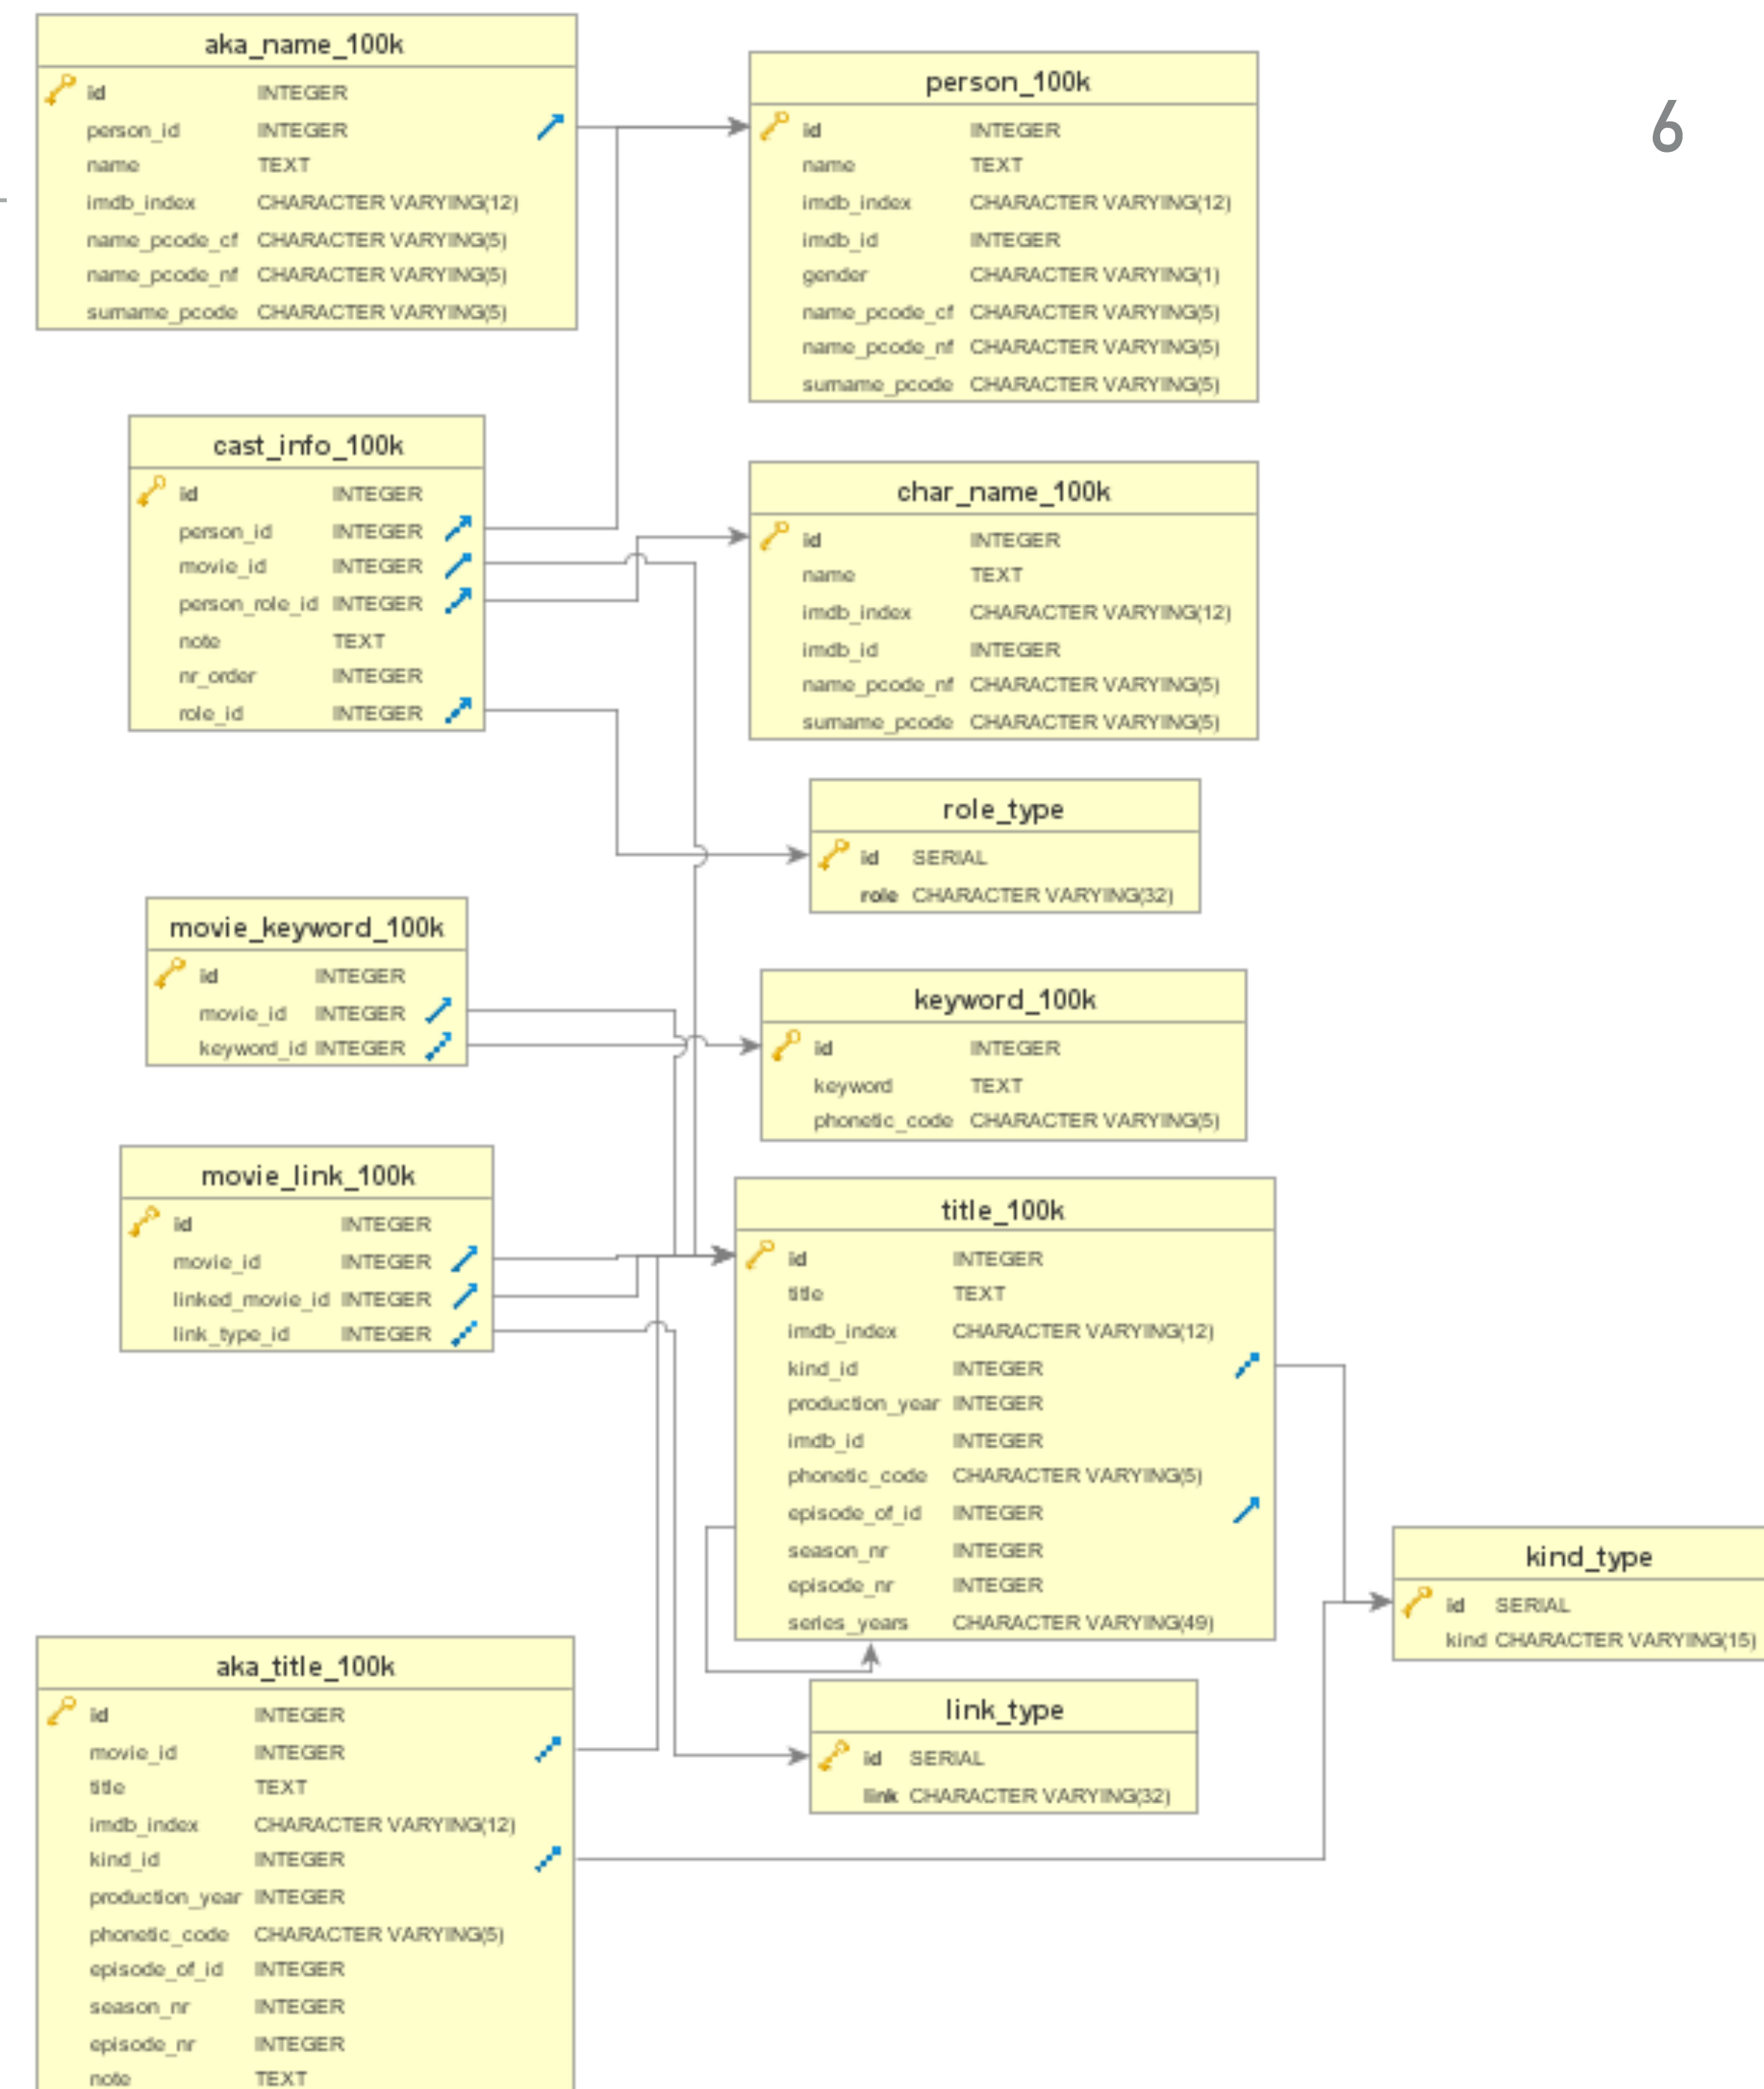
| Department | | |
|----------------|-----------------|----------|
| DeptName | Address | City |
| Administration | Bond Street | London |
| Production | Rue Victor Hugo | Toulouse |
| Distribution | Pond Road | Brighton |
| Planning | Bond Street | London |
| Research | Sunset Street | San Joné |

EXAMPLE DB2: PRODUCTS

| Supplier | | | | Supply | | | Products | | | | |
|--------------|--------|--------------|-----------|--------------|--------------|--------|--------------|---------|-------|------|------------|
| <u>CodeS</u> | NameS | Shareholders | Office | <u>CodeS</u> | <u>CodeP</u> | Amount | <u>CodeP</u> | NameP | Color | Size | Storehouse |
| S1 | John | 2 | Amsterdam | S1 | P1 | 300 | P1 | Sweater | Red | 40 | Amsterdam |
| S2 | Victor | 1 | Den Haag | S1 | P2 | 200 | P2 | Jeans | Green | 48 | Den Haag |
| S3 | Anna | 3 | Den Haag | S1 | P3 | 400 | P3 | Shirt | Blu | 48 | Rotterdam |
| S4 | Angela | 2 | Amsterdam | S1 | P4 | 200 | P4 | Shirt | Blu | 44 | Amsterdam |
| S5 | Paul | 3 | Utrecht | S1 | P5 | 100 | P5 | Skirt | Blu | 40 | Den Haag |
| | | | | S1 | P6 | 100 | P6 | Coat | Red | 42 | Amsterdam |
| | | | | S2 | P1 | 300 | | | | | |
| | | | | S2 | P2 | 400 | | | | | |
| | | | | S3 | P2 | 200 | | | | | |
| | | | | S4 | P3 | 200 | | | | | |
| | | | | S4 | P4 | 300 | | | | | |
| | | | | S4 | P5 | 400 | | | | | |

EXAMPLE DB3: IMDB

- ▶ A subset of the schema and data from the [IMDB.com](https://www.imdb.com) website
 - ▶ Actors (person_100k), Movies (title_100k), and Actors in Movies (cast_info_100k)
 - ▶ Plus aliases, keywords, movie genres, etc.
- ▶ We will use MongoDB and Neo4J implementations of the same database (obviously, with different schemas)
- ▶ Get it (with import instructions) here
 - ▶ https://docs.google.com/document/d/1jj3cMAnk6Rc0mHkkOAIYDzYLjKisCuyj4-3KF9l-_8o



WITH QUERIES

THE WITH CLAUSE

- ▶ A.k.a. **Common Table Expressions** (CTE)
- ▶ Allows the definition of a table (or multiple tables) that can be used only within a specific query
 - ▶ Similar to a VIEW (see later)
- ▶ But also the execution of INSERT/UPDATE/DELETE operations within the same query
 - ▶ RETURNING clauses give access to processed rows

EXAMPLES /1

- ▶ Retrieve pairs of titles from the 00s that have the same name of a title from the 80s


```
WITH NinetiesMovies AS(  
  SELECT title, production_year  
  FROM title_100k, kind_type  
  WHERE production_year BETWEEN 1990 AND 1999  
  AND title_100k.kind_id = kind_type.id  
  AND kind = 'movie'  
) , EightiesMovies AS (  
  SELECT title, production_year  
  FROM title_100k, kind_type  
  WHERE production_year BETWEEN 1980 AND 1989  
  AND title_100k.kind_id = kind_type.id  
  AND kind = 'movie'  
)  
  
SELECT *  
FROM NinetiesMovies, EightiesMovies  
WHERE NinetiesMovies.title = EightiesMovies.title
```

| | title text | production_year integer | title text | production_year integer |
|---|---------------|----------------------------|---------------|----------------------------|
| 1 | Going Down | 2005 | Going Down | 1988 |
| 2 | Mai | 2009 | Mai | 1989 |
| 3 | Run | 2009 | Run | 1989 |
| 4 | Urge | 2000 | Urge | 1988 |
| 5 | Willow | 2007 | Willow | 1988 |

EXAMPLES /2

- ▶ Remove all aliases containing the name *Alessandro*, but copy them in a log table

```
create TABLE aka_name_log(  
  like aka_name_100k  
  including all  
  including constraints  
  including indexes  
)
```



```
WITH moved_rows AS (  
  DELETE FROM aka_name_100k  
  WHERE name LIKE '%Alessandro%'  
  RETURNING *  
)  
INSERT INTO aka_name_log  
SELECT * FROM moved_rows
```

RECURSIVE QUERIES

- ▶ The WITH clause enables the definition of recursive queries
 - ▶ Useful to navigate recursive relationships

```
WITH RECURSIVE t(n) AS (  
  VALUES (1)  
  UNION ALL  
  SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n) FROM t
```

Base query

EXAMPLE

```
WITH RECURSIVE included_parts(sub_part, part, quantity) AS (  
    SELECT sub_part, part, quantity  
    FROM parts  
    WHERE part = 'our_product'  
    UNION ALL  
    SELECT p.sub_part, p.part, p.quantity  
    FROM included_parts pr, parts p  
    WHERE p.part = pr.sub_part  
)  
  
SELECT sub_part, SUM(quantity) as total_quantity  
FROM included_parts  
GROUP BY sub_part
```

VIEWS

VIEWS (VIRTUAL TABLES) IN SQL

Any relation that is not of the conceptual model but is made visible to a user

- ▶ Single table derived **from other tables**
- ▶ Considered to be a **virtual** relation
- ▶ Always up-to-date
 - ▶ "Physically" or "logically"
 - ▶ The DBMS takes care of synchronisation

External Level

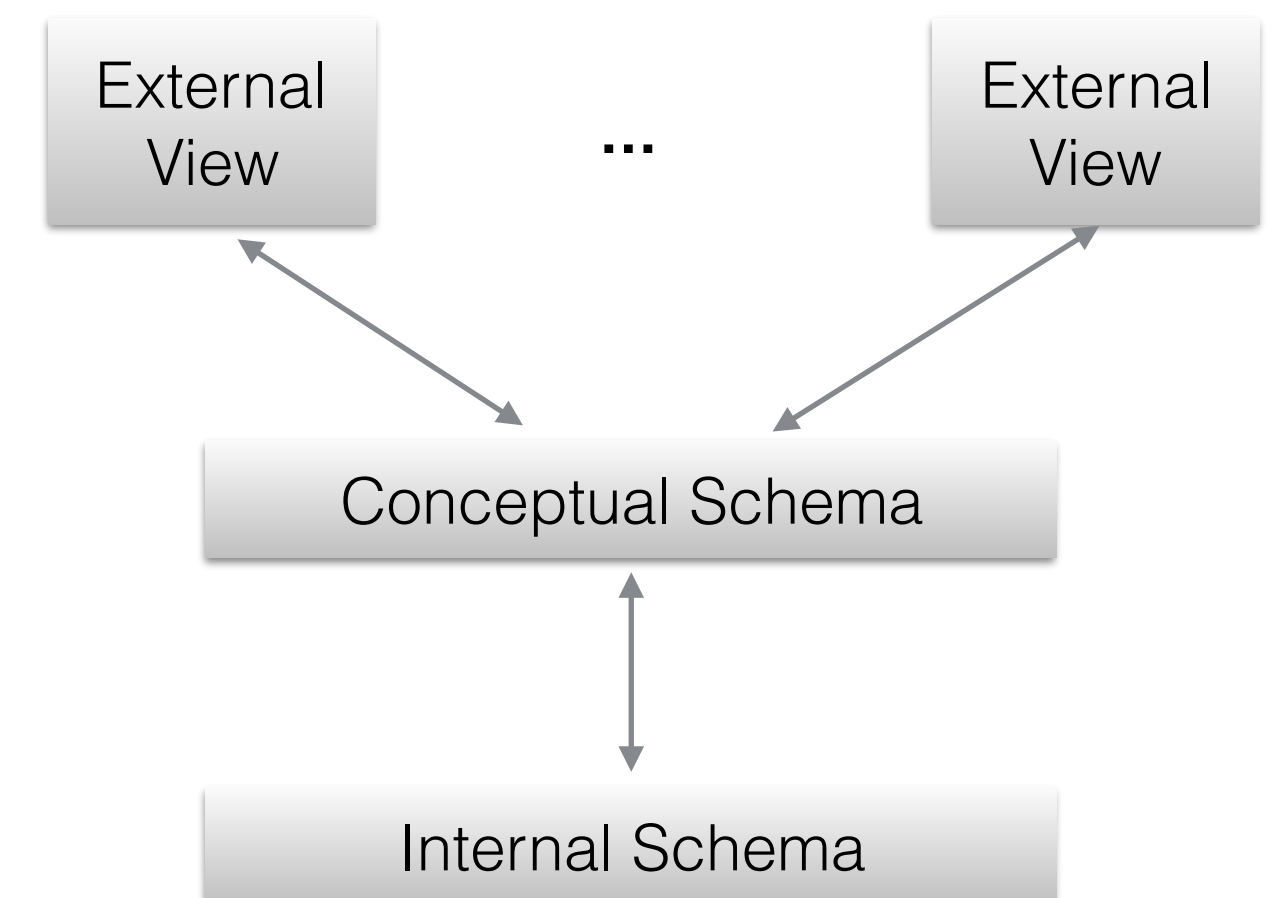
External/Conceptual
Mapping

Conceptual Level

Conceptual/Internal
Mapping

Internal Level

End Users



WHY VIEWS?

- ▶ Users **should not be allowed** to access the **actual** relations / data stored in the database
 - ▶ Limits on row values
 - ▶ Aliases for column or table names
- ▶ Favour the reuse of queries that are frequently used just to **denormalize** the logical schema
- ▶ Help simplifying complex queries, or enable the execution of queries that could not be expressed otherwise
 - ▶ **Decompose** the problem and produce a **more readable** solution
 - ▶ Combine and nest several **aggregate operators**

VIEWS DEFINITION

```
CREATE [TEMPORARY] VIEW ViewName [(AttributeList)]  
AS SelectSQL  
[WITH [CASCADED|LOCAL] CHECK OPTION]
```

- ▶ ViewName
 - ▶ Used by other queries / views as table name
- ▶ [(AttributeList)]
 - ▶ Used for projection / renaming / functions / operations
- ▶ SelectSQL
 - ▶ Any SELECT query (with some limitations)
- ▶ [WITH [CASCADED|LOCAL] CHECK OPTION]
 - ▶ Constrains inserts or updates to rows in tables referenced by the view
- ▶ [TEMPORARY]
 - ▶ Defines if the view should be dropped at the end of the current user session

VIEW DEFINITION EXAMPLES

► Administration Employees

```
CREATE VIEW AdminEmployee (FirstName, Surname, Salary) AS  
SELECT FirstName, Surname, Salary  
FROM Employee  
WHERE Dept = 'Administration' AND Salary > 10
```

VIEW

► Junior Administration Employees

```
CREATE VIEW JuniorAdminEmployee AS  
SELECT *  
FROM AdminEmployee  
WHERE Salary < 50  
WHICH CHECK OPTION
```

VIEW of VIEWS

USING VIEWS FOR QUERYING /1

- Find the departments with the highest salary expenditure

Without VIEWS

```
SELECT Dept
FROM Employee
GROUP BY Dept
HAVING sum(Salary) >= ALL(SELECT Dept
                           FROM Employee
                           GROUP BY Dept)
```

With VIEWS

```
CREATE VIEW SalaryBudget(Dept,SalaryTotal) AS
SELECT Dept, sum(Salary)
FROM Employee
GROUP BY Dept
```

```
SELECT Dept
FROM SalaryBudget
WHERE SalaryTotal =(SELECT max(SalaryTotal)
                    FROM SalaryBudget)
```

USING VIEWS FOR QUERYING /2

- Find the average number of offices per department

Incorrect Solution

```
SELECT avg(count(DISTINCT Office))  
FROM Employee  
GROUP BY Dept
```

SQL does not allow cascades
of aggregate operators

Correct Solution

```
CREATE VIEW DeptOffice(Dept,NoOfOffice) AS  
SELECT Dept, count(DISTINCT Office)  
FROM Employee  
GROUP BY Dept
```

```
SELECT avg(NoOfOffices)  
FROM DeptOff
```

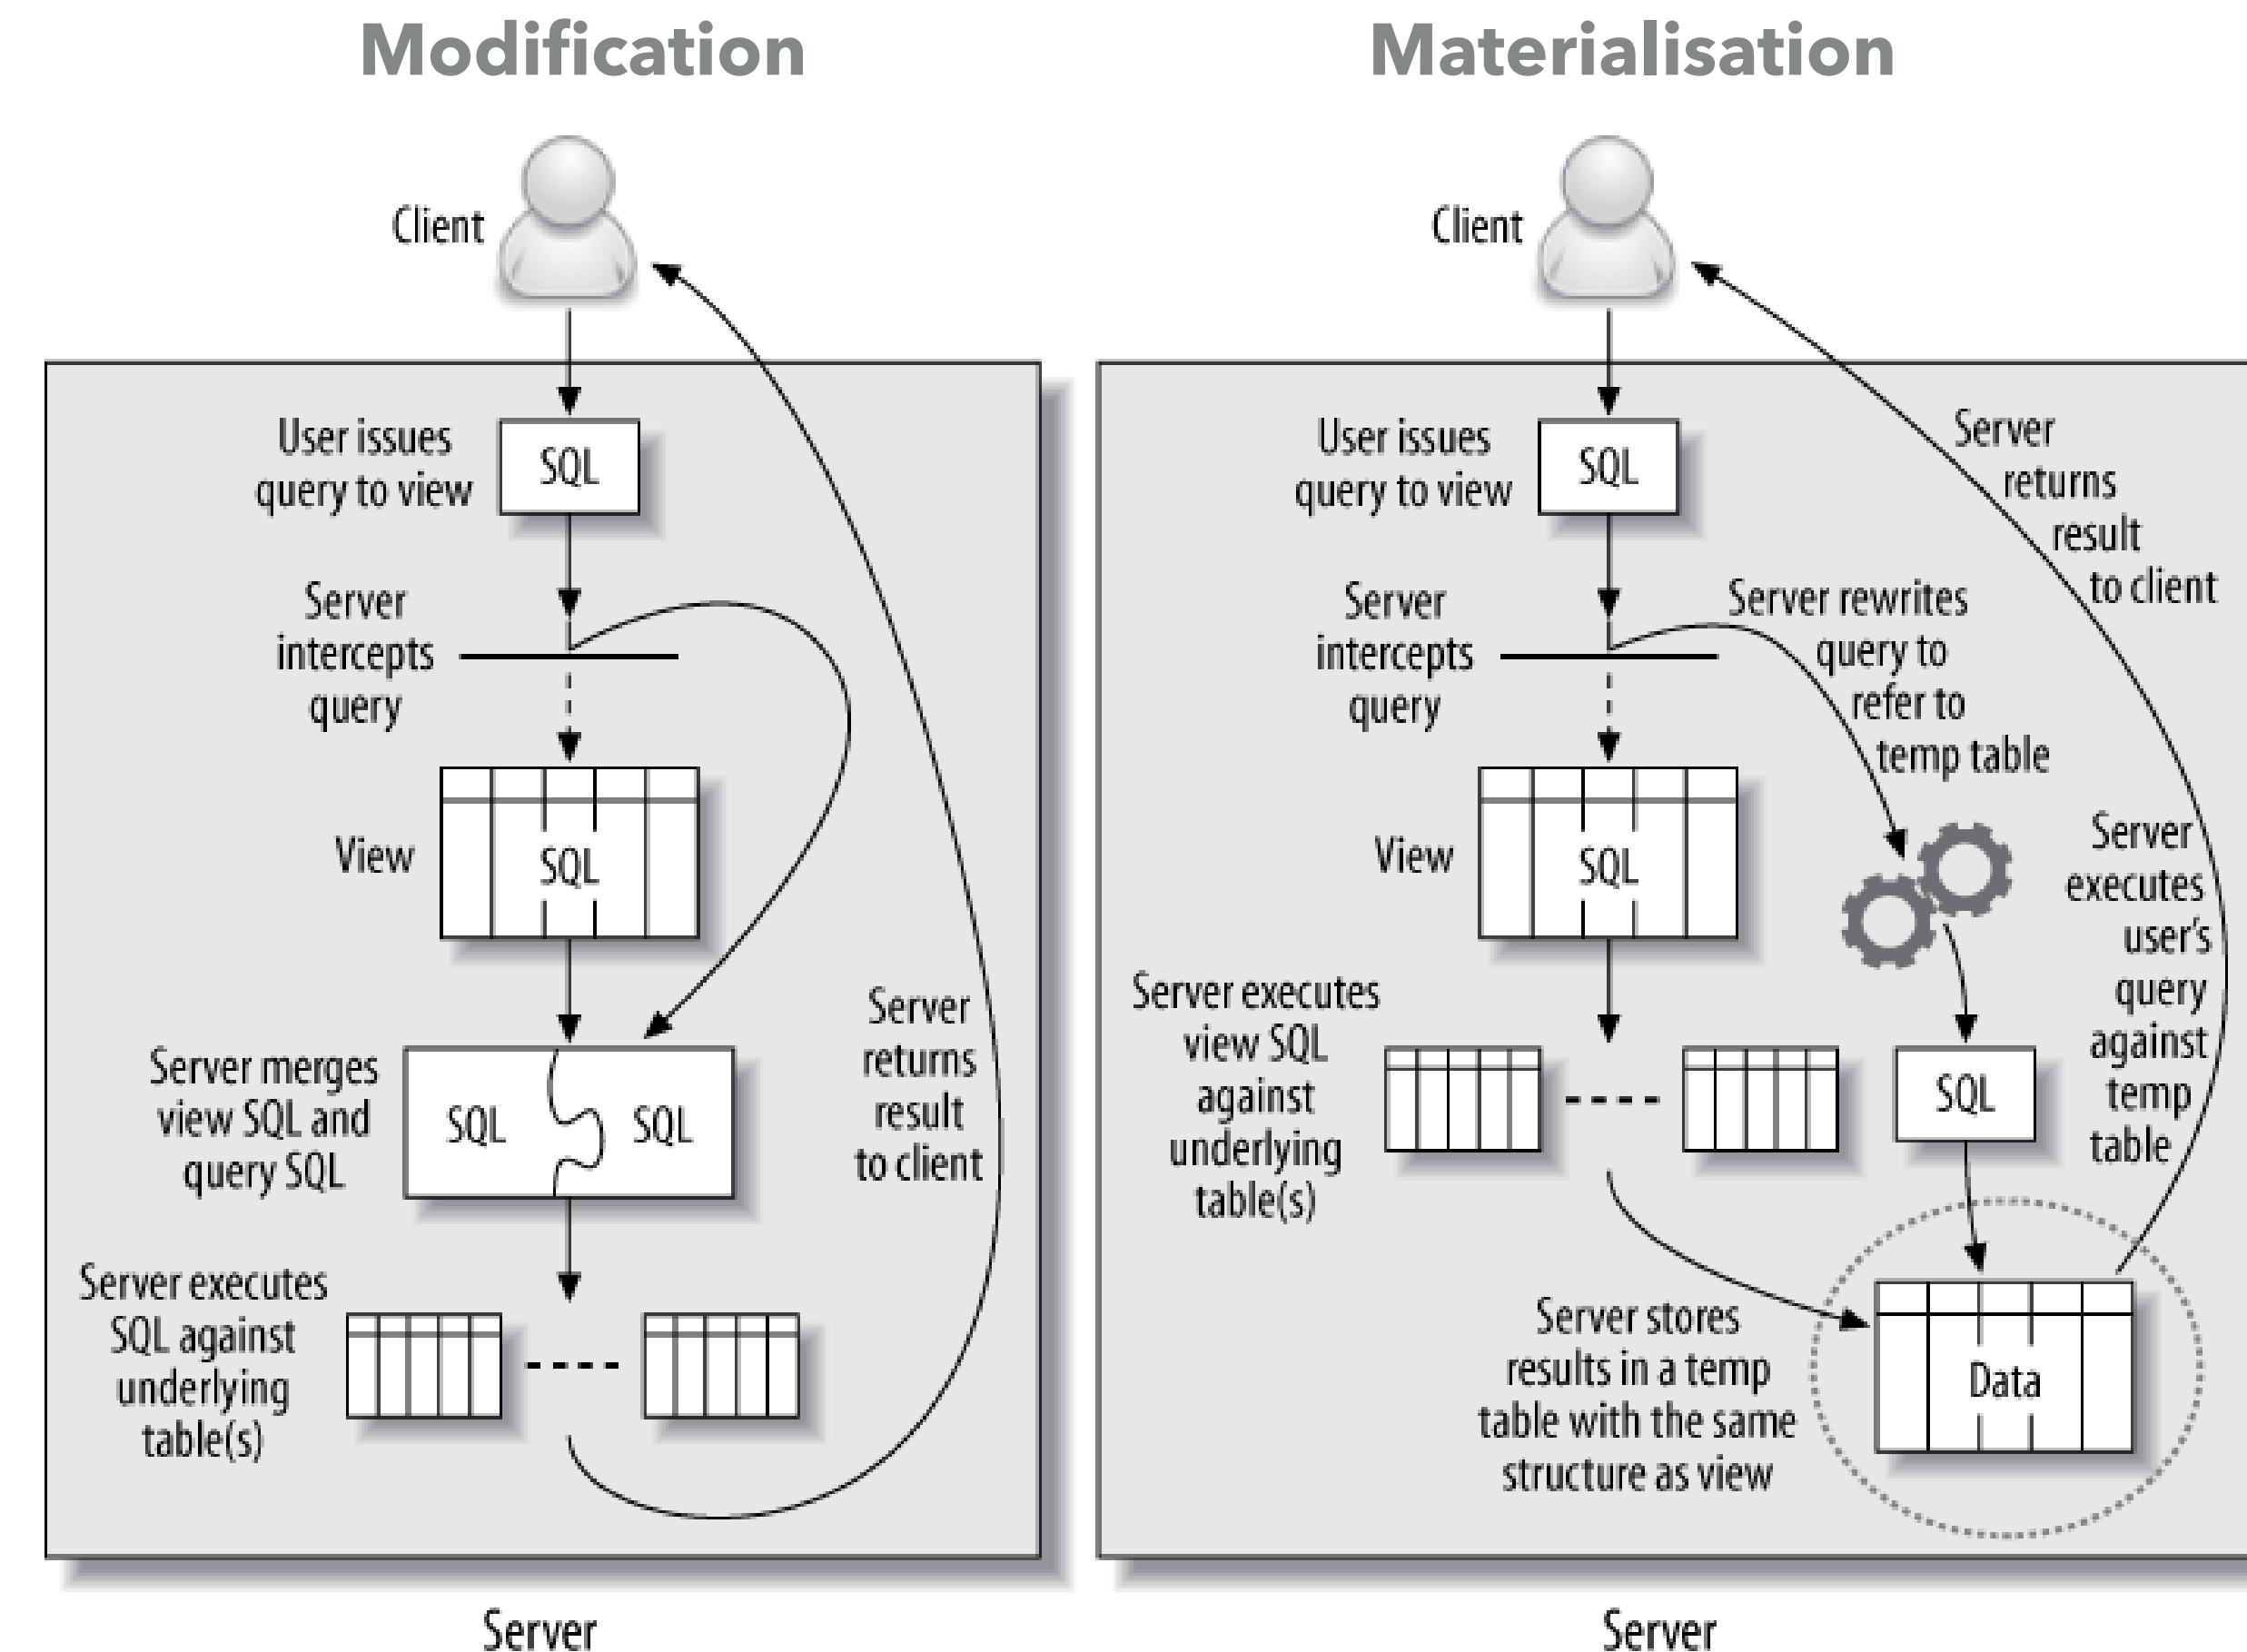
VIEW EXECUTION APPROACHES

► Query Modification

- **Modify** view query into a query on underlying base tables
- Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute

► View Materialisation

- **Physically create** a temporary view table when the view is first queried
- **Keep that table** on the assumption that other queries on the view will follow
- Requires efficient strategy for **automatically** updating the view table when the base tables are updated



Source: High Performance MySQL, 3rd Edition. TEMPTABLE mechanism in MySQL works similarly to POSTGRESQL Materialised View extension, In the latter, materialised views are updated on demand

VIEW EXECUTION IN POSTGRESQL

▶ Query Modification

- ▶ The query is run every time the view is referenced in a query
- ▶ Standard mode
- ▶ TEMPORARY option specify lifetime of view (current session)

▶ Query Materialisation: CREATE MATERIALIZED VIEW

- ▶ Defines a physical table
- ▶ Manual refresh of data
 - ▶ REFRESH MATERIALIZED VIEW
 - ▶ Can slow down/lock the DB
- ▶ POSTGRESQL Extension

```
CREATE MATERIALIZED VIEW [CONCURRENTLY] ViewName  
AS SelectSQL  
[WITH [NO] DATA]
```

AUTOMATICALLY UPDATABLE VIEWS

Idea: Update the underlying base tables through CRUD operations on a VIEW

- ▶ INSERT, UPDATE or DELETE statement on the view are turned into corresponding statements on the base relation
- ▶ Possible when
 - ▶ VIEW defined on single table / VIEW
 - ▶ No WITH, DISTINCT, GROUP BY, HAVING, LIMIT, or OFFSET clauses at the top level
 - ▶ No set operations (UNION, INTERSECT or EXCEPT) at the top level
 - ▶ No aggregates or set-returning functions

Some updates can be ambiguous

```
CREATE VIEW Employee_info AS  
  SELECT FirstName, Surname, City  
  FROM Employee JOIN Department ON dept=DeptName
```

- ▶ Which department, if multiple departments in Berlin?
- ▶ What if no department is in Berlin?

```
INSERT INTO Employee_info VALUES ('Joe', 'Smith', 'Berlin');
```

AUTOMATICALLY UPDATABLE VIEWS

Idea: Update the underlying base tables through CRUD operations on a VIEW

- ▶ There might be a mix of updatable and non-updatable column
 - ▶ Updatable: simple reference to updatable column of the base relation
 - ▶ Non-updatable: error when INSERT or UPDATE
- ▶ If WHERE statement in the query
 - ▶ Only rows that are in the view can be deleted or updated
 - ▶ INSERT can be executed on the view, but rows might not be in it
- ▶ CHECK OPTION
 - ▶ Ensures that any rows changed through the view continue to match the view's WHERE clause after the changes
 - ▶ LOCAL: check only the conditions defined directly in the view
 - ▶ CASCADE: check also conditions of underlying base VIEWS

LIMITATIONS AND PERFORMANCE IMPLICATIONS

- ▶ User, system, or local variables are **not allowed** in the SQL SELECT statement
 - ▶ It is not possible to have parametric VIEWS
- ▶ If base table schema changes, VIEWS won't be valid anymore
 - ▶ There is no VIEW integrity constraint
- ▶ VIEWS can have bad performance
 - ▶ But sometimes better than an equivalent query that doesn't use a view (caching)
- ▶ MATERIALIZED views add overhead, and it's hard to predict how a view will impact performance
 - ▶ VIEWS of VIEWS can easily generate VERY complex execution plans

VIEWS VS. BASE TABLES

BASE TABLE

- ▶ Tuples always physically stored in database
- ▶ CRUD operations are always allowed
- ▶ Always possible to **define tables** of views

VIEW

- ▶ Tuples do not necessarily exist in physical form
- ▶ CRUD operations are not always allowed on views
- ▶ Can be used to create other views, but not in a mutually dependent way
 - ▶ Recursion is possible on same VIEW

VIEWS IN MONGODB AND NEO4J

- ▶ MongoDB has a `createView` command
 - ▶ Same function as SQL VIEWS
 - ▶ Read only
 - ▶ Also VIEWS of VIEWS
- ▶ Neo4J has no VIEWS

ASSERTIONS AND TRIGGERS

SEMANTIC / BUSINESS CONSTRAINTS

Constraints that cannot be directly expressed in the schemas of the data model

- ▶ Examples
 - ▶ *The salary of an employee should not exceed the salary of the employee's supervisor*
 - ▶ *The maximum number of hours an employee can work on all projects per week is 56*
- ▶ SQL provides two constructs (not supported by all systems)
 - ▶ CREATE ASSERTION
 - ▶ CREATE TRIGGER

ASSERTIONS

- ▶ Assertions permit the specification of constraints outside of table definitions
 - ▶ Useful in many situations (e.g., to express generic inter- relational constraints)
- ▶ An assertion associates a name to a check clause

```
CREATE ASSERTION assertion_name CHECK (condition)
```

The constraint is satisfied if *no combination of tuples* in that database violates it

- ▶ The condition **must hold TRUE** for **every database state** for the assertion to be satisfied
- ▶ The DBMS is responsible for the condition not to be violated

ASSERTION EXAMPLE

- ▶ The condition clause can contain any condition that can be specified in the WHERE clause of a SELECT query
- ▶ There must always be at least one tuple in table EMPLOYEE

```
CREATE ASSERTION AlwaysOneEmployee  
  CHECK (1 <= (SELECT count(*)  
               FROM Employee)  
        )
```

Not available in PostgreSQL

TRIGGERS

Triggers (Active Rules): rules that are automatically triggered by database **events**, and that initiate certain **action** if certain **conditions** are met

- ▶ A trigger has three logical components
- ▶ The triggering **event**
 - ▶ data operations, temporal events, external events
- ▶ The **condition** that determines whether the rule action should be executed
 - ▶ Optional
- ▶ The **action** to be taken
 - ▶ SQL statements, external programs, etc.

WHAT ARE TRIGGERS FOR?

- ▶ **Notify** users about **violations** of some constraints
 - ▶ E.g. a manager should be notified if an employee is having too much travel expenses
- ▶ Manage **advanced referential integrity** constraints not based on keys
 - ▶ E.g. avoid updates of tuple values based on values of other tuples
- ▶ Automatic **maintenance** of **derived** data
 - ▶ E.g. materialised views, data replication, etc.

GRANULARITY AND EXECUTION MODE

- ▶ Granularity
 - ▶ **Row-level:** the trigger is activated once for every tuple on which the event occurred
 - ▶ **Statement-level:** the trigger is activated once for every SQL statement, referring to all the tuples on which the statement operated (set-oriented)
- ▶ Execution mode
 - ▶ **Immediate:** right after (or even before) the event
 - ▶ **Before:** to perform actions prior to changes in the table. New /modified record can be changed
 - ▶ **After:** to perform actions after changes in the table. Record written in the table cannot be changed
 - ▶ **Deferred:** at transaction commit

TRIGGERS IN POSTGRESQL

Examples in book use ORACLE syntax
PostgreSQL complete syntax has more options

```
CREATE TRIGGER name {BEFORE|AFTER}{event [OR ...]}  
ON tableName  
[FOR [EACH] {ROW|STATEMENT}]  
[WHEN (condition)]  
EXECUTE PROCEDURE function_name (arguments)
```

name

- ▶ Unique for the table

time

- ▶ BEFORE
- ▶ AFTER

event

- ▶ INSERT/UPDATE/DELETE
- ▶ Multiple events can be specified using OR (a PostgreSQL extension of the standard)

function_name(arguments)

- ▶ The function to be executed (and input parameters)

FOR EACH ROW

- ▶ function_name is executed once for each of the affected records

FOR EACH STATEMENT

- ▶ function_name is executed once for any given operations

STORED PROCEDURES IN POSTGRESQL

```
CREATE FUNCTION name (arguments)
RETURNS returnType
local_declarations
function_body
```

- ▶ Programs stored in the database
 - ▶ Can be written in different imperative/declarative languages
 - ▶ PL/pgSQL is PostgreSQL language
- ▶ A trigger function
 - ▶ Takes no parameters
 - ▶ Return type is TRIGGER
- ▶ **OLD** and **NEW** keywords: allow to refer to the data before and after the activating event takes place
- ▶ It is possible to use variables and flow controls.

TRIGGER EXAMPLE (FROM WERKCOLLEGE 5)

- ▶ Automatically increase/decrease vote counts when votes are created/deleted

```
CREATE FUNCTION f_inc_votes()
RETURNS TRIGGER
LANGUAGE PLPGSQL;
AS $$
BEGIN
    UPDATE suggestions
    SET votecount = votecount+1
    WHERE id = NEW.suggestionid;
    RETURN NEW;
END
$$

CREATE TRIGGER inc_votes
            AFTER INSERT ON votes
    FOR EACH ROW EXECUTE PROCEDURE
                        f_inc_votes();
```

```
CREATE FUNCTION f_dec_votes()
RETURNS TRIGGER
LANGUAGE PLPGSQL;
AS $$
BEGIN
    UPDATE suggestions
    SET votecount = votecount-1
    WHERE id = NEW.suggestionid;
    RETURN OLD;
END
$$

CREATE TRIGGER dec_votes
            AFTER DELETE ON votes
    FOR EACH ROW EXECUTE PROCEDURE
                        f_dec_votes();
```

EXTENSIONS (NOT USUALLY AVAILABLE)

- ▶ Boolean combinations of events
 - ▶ PostgreSQL has OR
- ▶ `instead of` clause
 - ▶ it is not executed the operation that triggered the event, but another one in its place
 - ▶ Available in PostgreSQL
- ▶ “Detached” execution mode: a separate transaction is started to manage the triggers
- ▶ Explicit user-defined priorities
- ▶ Rule sets, that can be made activated and deactivated with a single command

PROPERTIES OF TRIGGERS

► Termination


- A rule set is guaranteed to terminate if, for any database state and initial modification, rule processing cannot continue forever

► Confluence

- A rule set is confluent if, for any database state and initial modification, the final database state after rule processing is unique, i.e. it is independent of the order in which activated rules are executed

► Termination is assessed studying rules interaction

- An important conceptual tool is the *triggering graph*



ADVANTAGES OF TRIGGERS

- ▶ Provide a **complementary**, and **more robust**, integrity checking mechanism to foreign keys
 - ▶ They can verify if foreign key **tuples** have certain characteristics
- ▶ Ability to **catch** business process **errors**
- ▶ Guarantee that for **every** change, the trigger is run
 - ▶ Control code in external application could not control other 3rd party changes
- ▶ Allow **scheduled tasks** directly in the DB
 - ▶ No cron or other system-level scheduler
- ▶ Simplify **propagation** of value **changes** from various tables
- ▶ Enable the execution of **calculation before** a row or value is inserted

DISADVANTAGES OF TRIGGERS

- ▶ They can't replace ALL types of validation
 - ▶ No substitute for client-side validation
- ▶ **Hard to create**, debug and maintain
 - ▶ Few developer tools
 - ▶ It is VERY easy to create cyclical rules
 - ▶ **termination** and **confluence**
 - ▶ Bugs are difficult to locate in trigger chains
- ▶ Very heterogeneous functionalities and support across vendors

TRIGGERS IN MONGODB AND NEO4J

- ▶ MongoDB *Stitch* triggers allow for database events handling
 - ▶ Events: CRUD on documents, Authentication
 - ▶ Actions: Javascript functions
- ▶ Neo4J has no specific commands for triggers
 - ▶ a “TransactionEventHandler” API that could be used to implement custom Java functions that are executed when transactions are about to be committed

ACCESS CONTROL

ACCESS CONTROL

- ▶ Every component of the schema can be protected (tables, attributes, views, etc.)

The owner of a resource (the creator) assigns privileges to the other users

- ▶ A predefined **role** (`postgres`, or the name of the operating system user that initialised the cluster) represents the database administrator and has complete access to all the resources
- ▶ A privilege is characterised by:
 - ▶ The **resource**
 - ▶ The user who **grants** the privilege
 - ▶ The user who **receives** the privilege
 - ▶ The **action** that is allowed on the resource
 - ▶ Whether or not the privilege can be **passed on** to other users

MANAGING ROLES

- ▶ Role: a database user, or a group of database users
 - ▶ The `pg_roles` catalog contains the current list of roles
- ▶ The role determines the access privileges
 - ▶ Login
 - ▶ Database creation
 - ▶ Role creation
- ▶ Group roles can be used to grant and revoke privileges to multiple users
 - ▶ Group roles have no LOGIN access privilege

```
CREATE ROLE name
```

```
DROP ROLE name
```

```
CREATE ROLE role_name LOGIN
```

```
GRANT group_role TO role_name
```

POSTGRES ROLE AUTHENTICATION

- ▶ Trust
 - ▶ Any role name in the roles list can access
 - ▶ Not recommended for multi-user or networked machines
- ▶ Password-based
 - ▶ PostgreSQL database passwords are separate from operating system user passwords
- ▶ SSL Certificates
- ▶ Plus a number of external authentication protocols and systems
 - ▶ GSSAPI, SSPI, LDAP, RADIUS

VIEWS IN MONGODB AND NEO4J

▶ MongoDB

- ▶ Role-based access control, similar to PostgreSQL
- ▶ Privileges are granted at database and collection level
- ▶ MongoDB starts with no authentication methods enabled
 - ▶ It supports internal, certificate-based, and protocol-based authentication

▶ Neo4J

- ▶ Authentication and authorisation available only in the enterprise edition
- ▶ Role-based access control
- ▶ Privileges granted on the whole graph
- ▶ Authentication with native, LDAP, and other authorisation provides

WRAPPING UP

TODAY WE COVERED

- ▶ The WITH clause
- ▶ Views
- ▶ Triggers
- ▶ Access Control Mechanism

END OF LECTURE