

CSE1500 – WEB AND DATABASE TECHNOLOGY

DB LECTURE 6

NOSQL / NEO4J

Alessandro Bozzon

cse1500-ewi@tudelft.nl

AT THE END OF THIS LECTURE, YOU SHOULD BE ABLE TO....

- ▶ **Describe** and **design** data-driven applications using Neo4J

NEO4J

WHY?

Relational Databases
(incredibly!)

are not good in managing relationships!

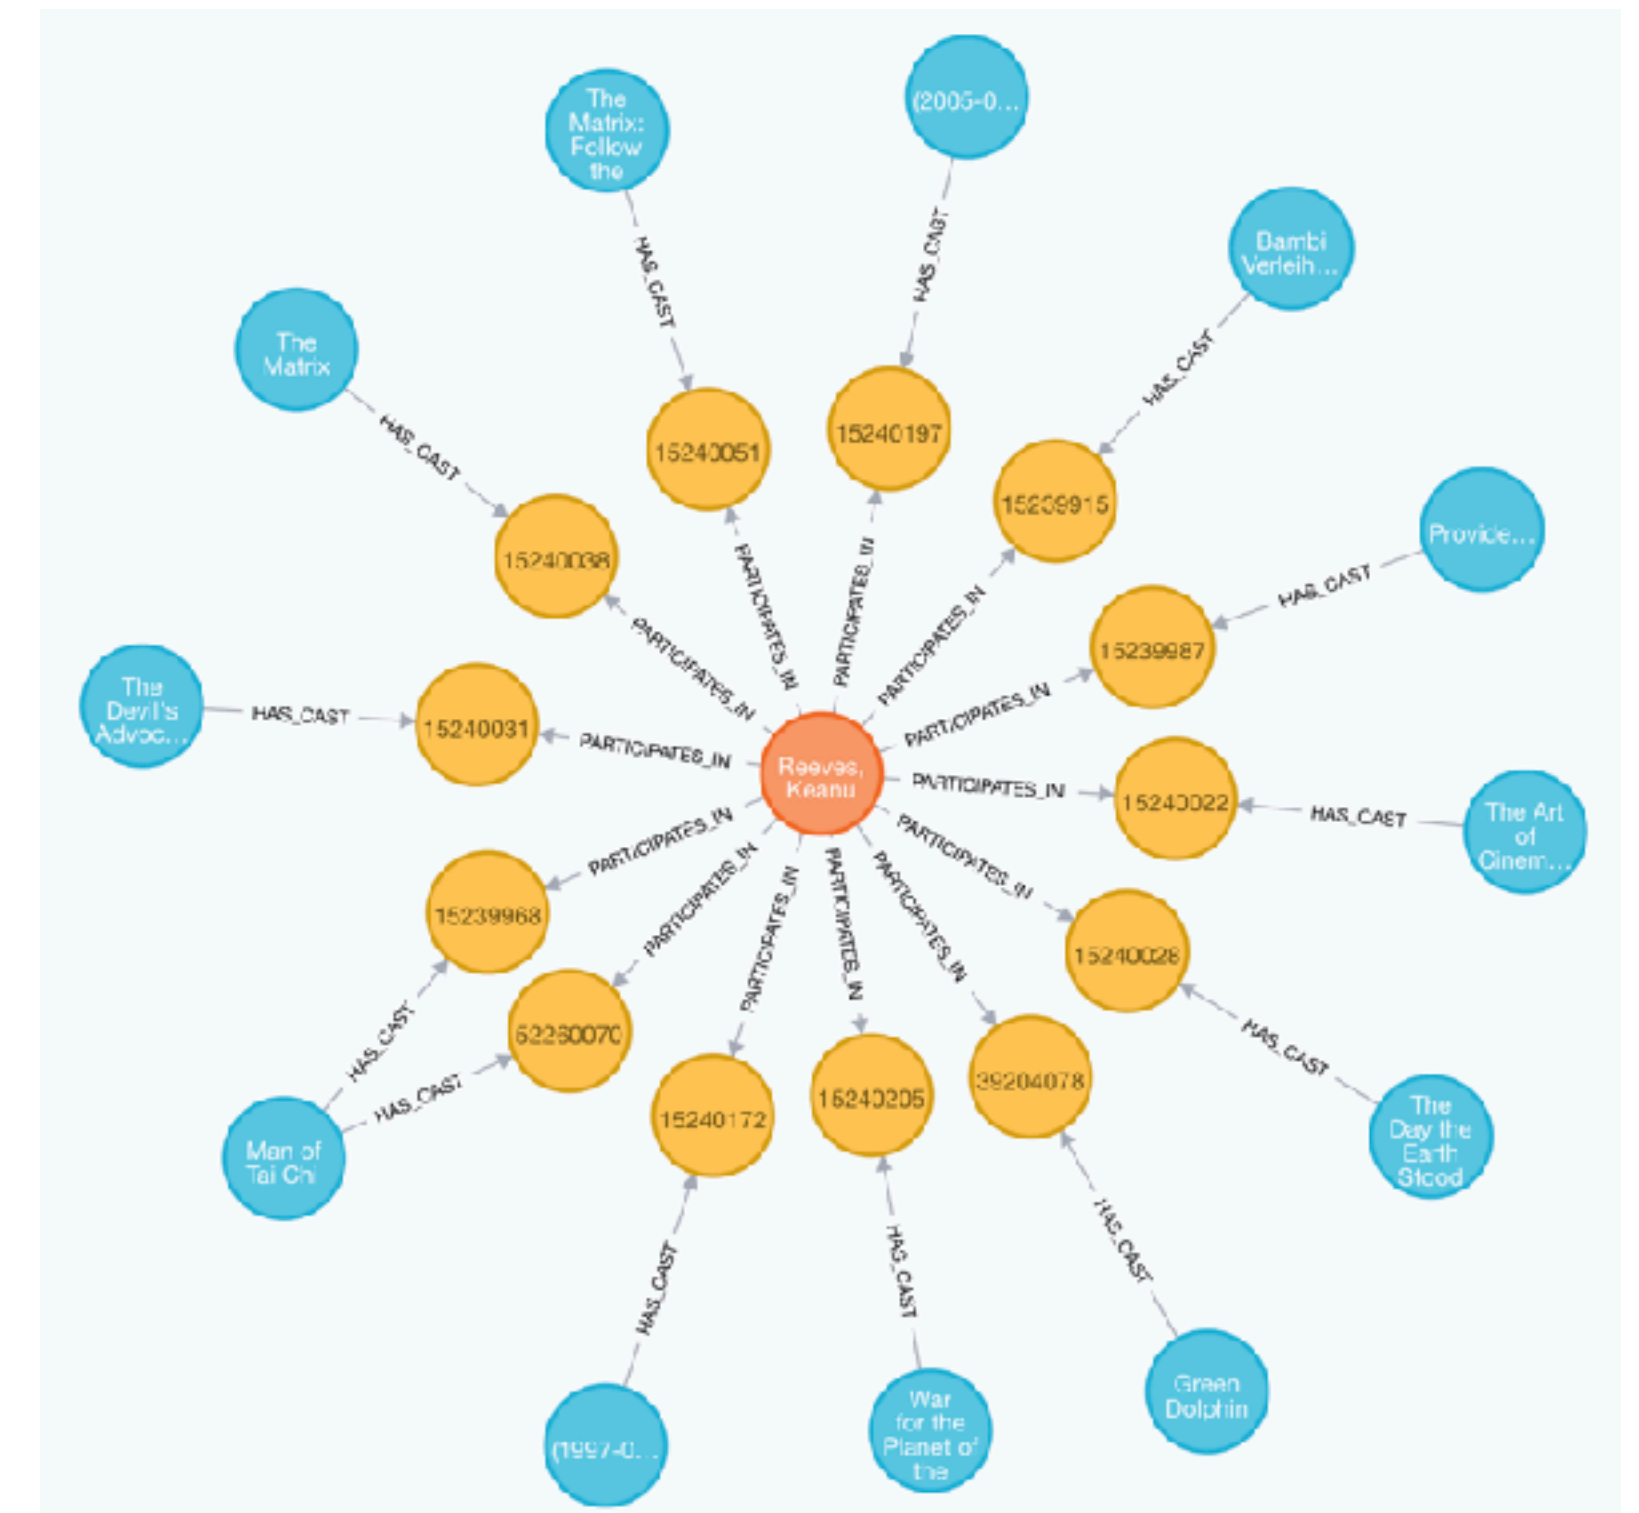

THE NEO4J GRAPH DATABASE

- ▶ Open source: <http://neo4j.com>
 - ▶ But with commercial – enterprise – edition (with caching, clustering, etc.)
 - ▶ Great documentation
- ▶ Implemented in Java, cross-platform
- ▶ Schemaless and “typeless”
 - ▶ Basic data types (e.g. string) , but no restrictions on entity and relationship types
- ▶ Declarative query language (**Cypher**)
 - ▶ On the path to standardisation <http://www.opencypher.org/>
- ▶ No sharding
- ▶ Master-Slave replication (but of entire graphs)
- ▶ **ACID** properties!



PROS/CONS OF GRAPH DATABASES

- Break data into even smaller units than RDBMs



- Not suitable for distribution, but ACID properties can be guaranteed (at a cost)
- Suitable for complex, semi-structured, highly connected data

APPLICATIONS

- ▶ Connected Data
 - ▶ Social graphs, knowledge graphs, citations
- ▶ Routing, Dispatch, Location-Based services
 - ▶ Every location is a node, connections between locations can be a (weighted) edge
- ▶ Recommendation engines
 - ▶ Establish relationships between people and things
- ▶ Don't use for massive data updates
 - ▶ It takes time to build and massively change graphs



GRAPH VS. RELATIONAL DATABASES

- ▶ RDBMs and NOSQL DBs lack relationships
 - ▶ JOINS and aggregates are sub-optimal solutions
 - ▶ Relationships (through reference or embedding) are not reflective
- ▶ Graph Databases embrace relationships
 - ▶ Traversal operation are highly efficient

**THIS IS WITHOUT RECURSIVE
COMMON TABLE EXPRESSIONS**



Friends of friends of my friends?


Depth	RDBMS execution time (s)	Neo4j execution time (s)	Records returned
2	0.016	0.01	~2500
3	30.267	0.168	~110,000
4	1543.505	1.359	~600,000
5	Unfinished	2.132	~800,000

- ▶ In a depth two (friend-of-friend), both RDBMs and Neo4J perform well enough
 - ▶ But when we do the depth three it clear that RDBMs can no longer keep up

NEO4J DATA MODEL

WHAT IS A GRAPH?

- ▶ Informally, a graph is a set of nodes joined by a set of directed or undirected edges
- ▶ A subgraph of a graph G is a graph fully contained by G



- ▶ Almost anything can be modelled as a graph

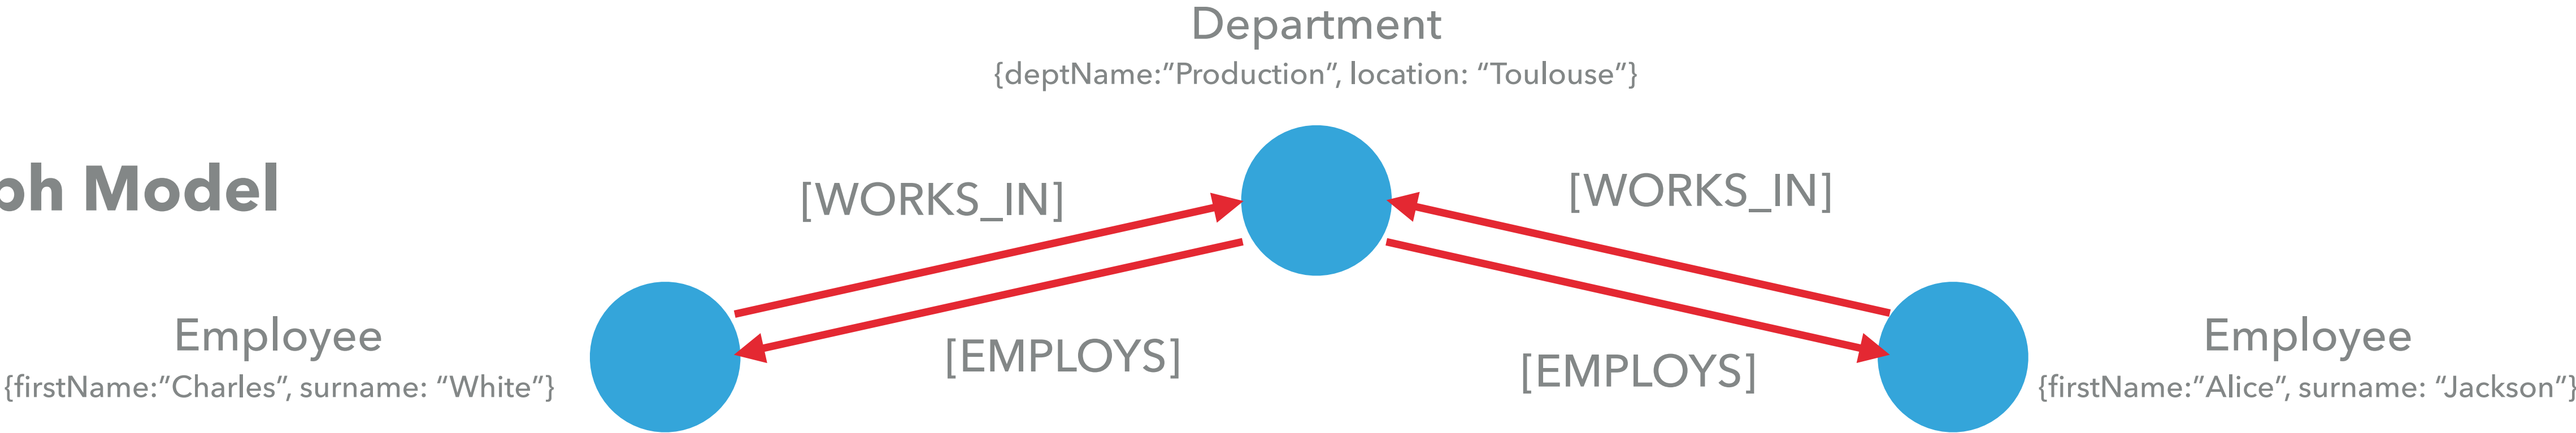
RELATIONAL DATABASE REPRESENTATION

Relational Model

Employee					
FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	White	Production	20	36	Toulouse
Gus	Green	Administration	20	40	Oxford
Jackson	Neri	Distribution	16	45	Dover
Charles	Brown	Planning	14	80	London
Laurence	Chen	Planning	7	73	Worthing
Pauline	Bradshaw	Administration	75	40	Brighton
Alice	Jackson	Production	20	46	Toulouse

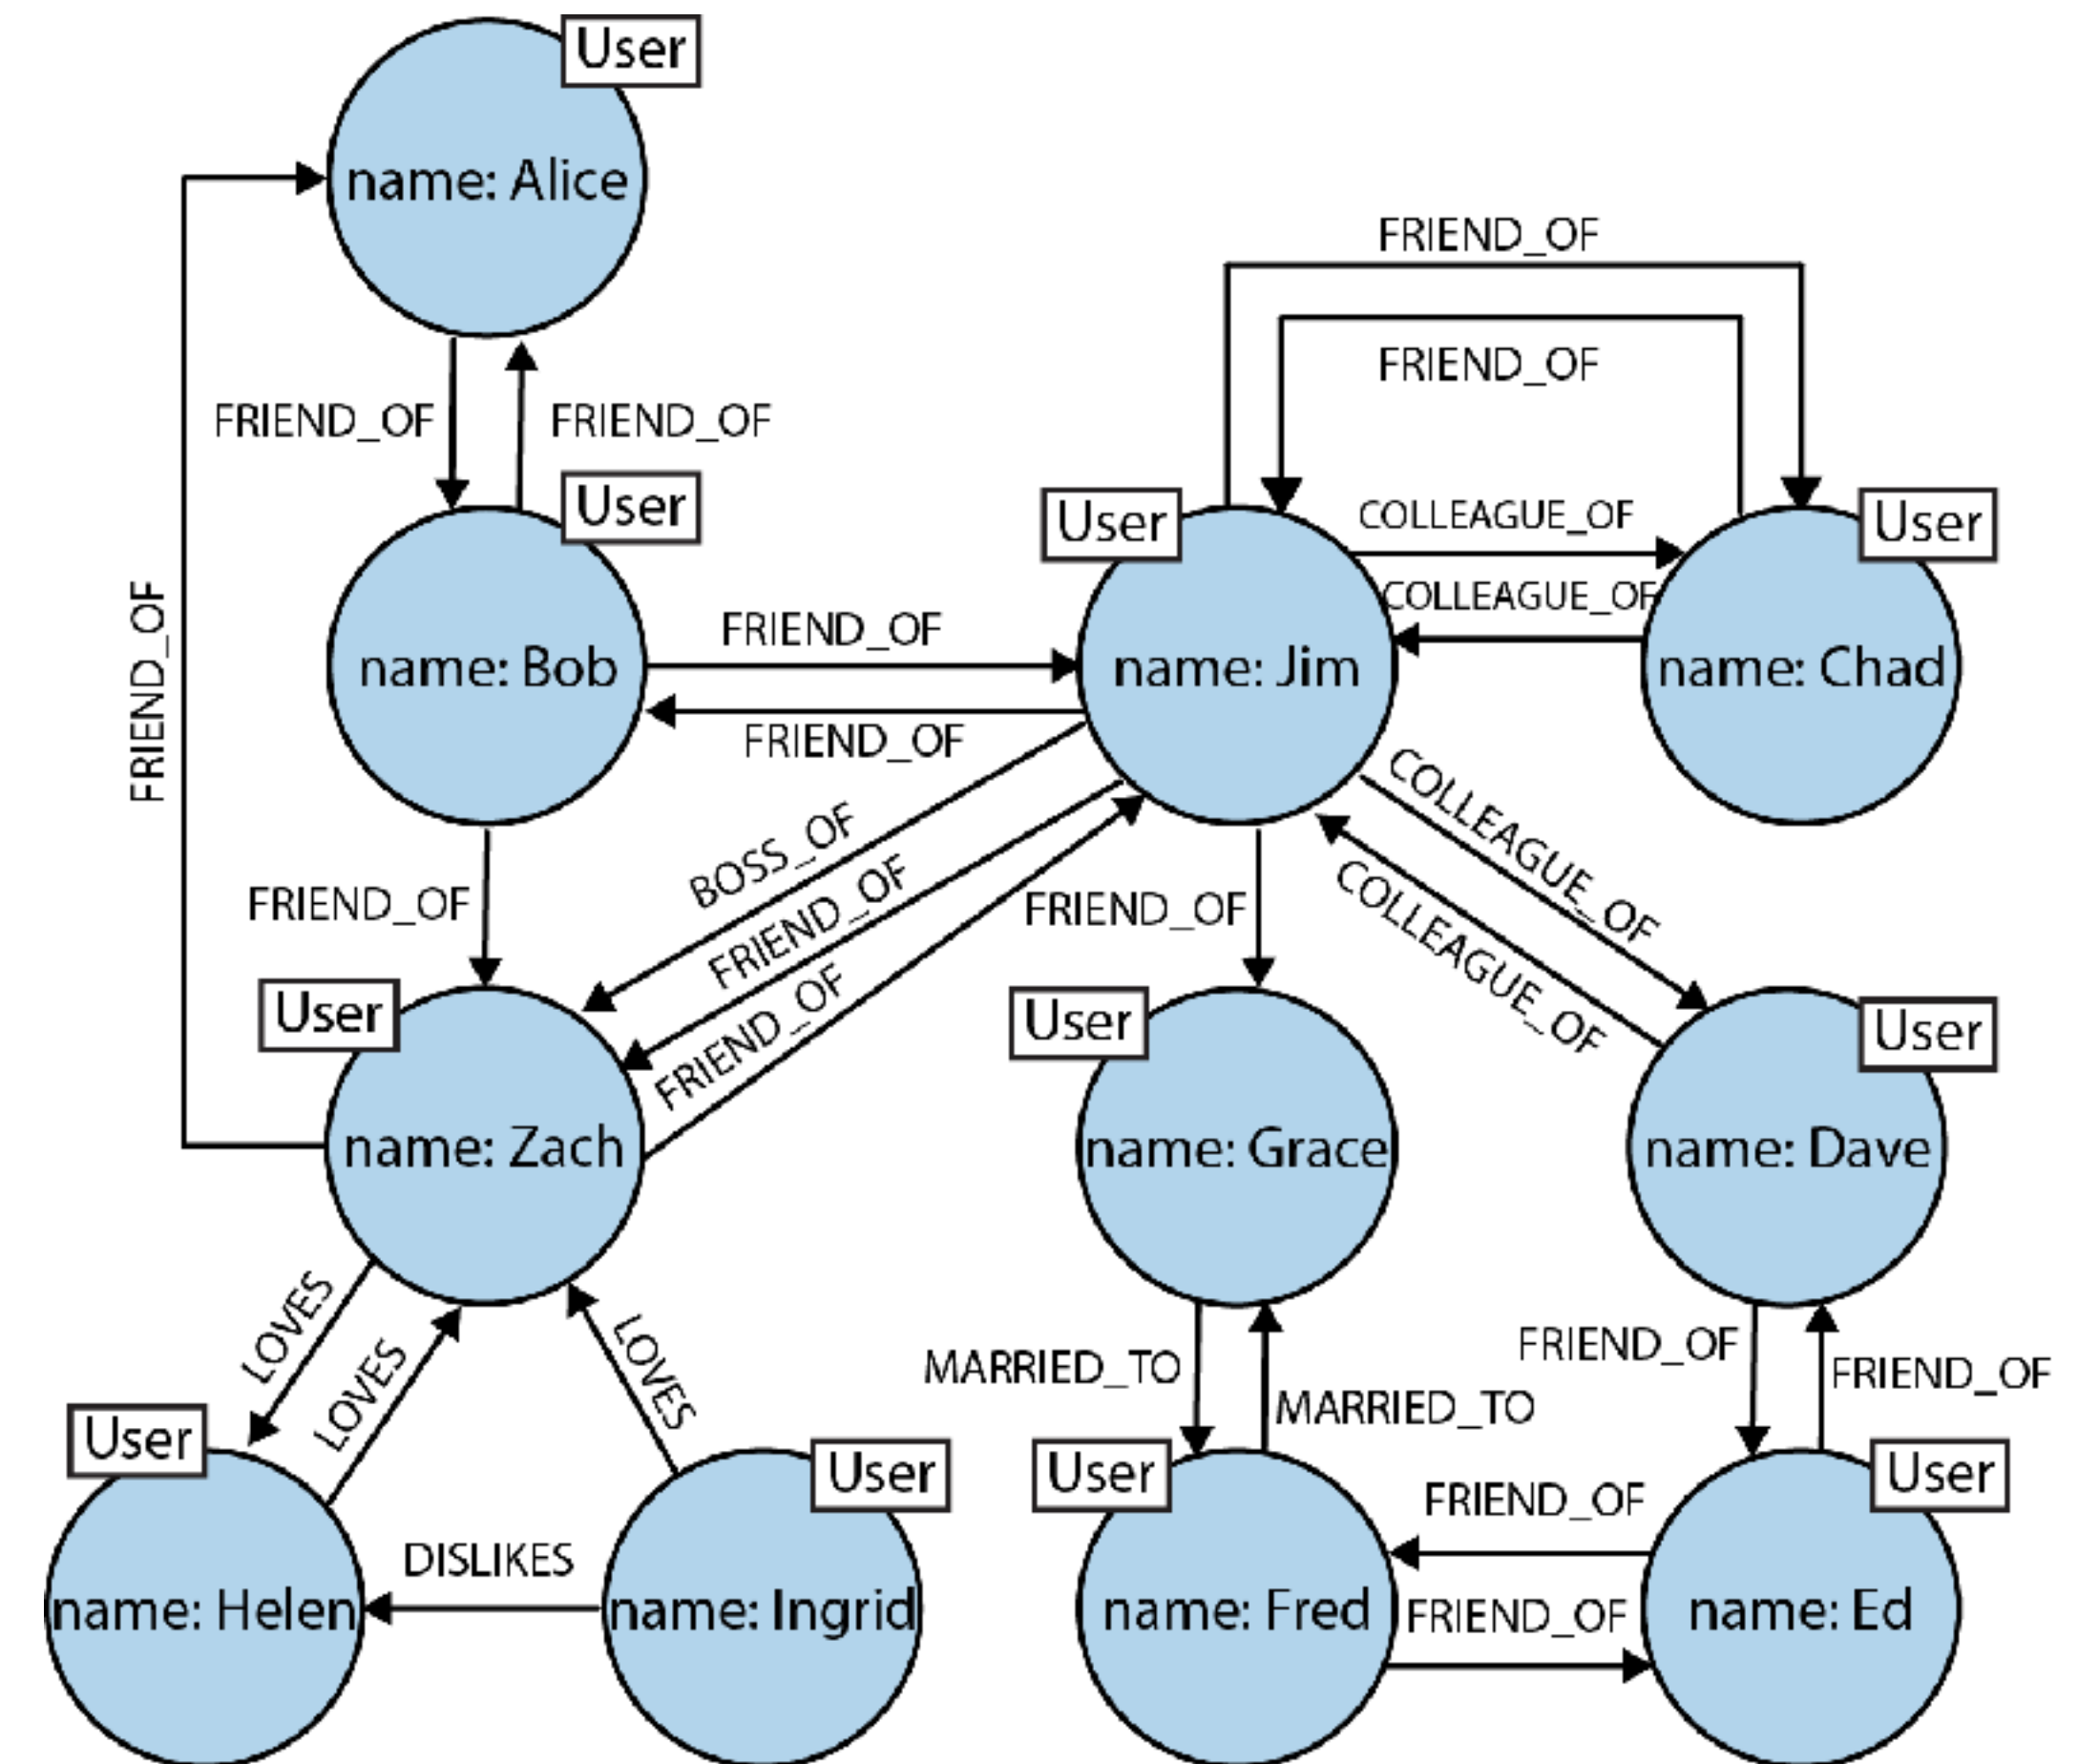
Department		
DeptName	Address	City
Administration	Bond Street	London
Production	Rue Victor Hugo	Toulouse
Distribution	Pond Road	Brighton
Planning	Bond Street	London
Research	Sunset Street	San Joné

Graph Model



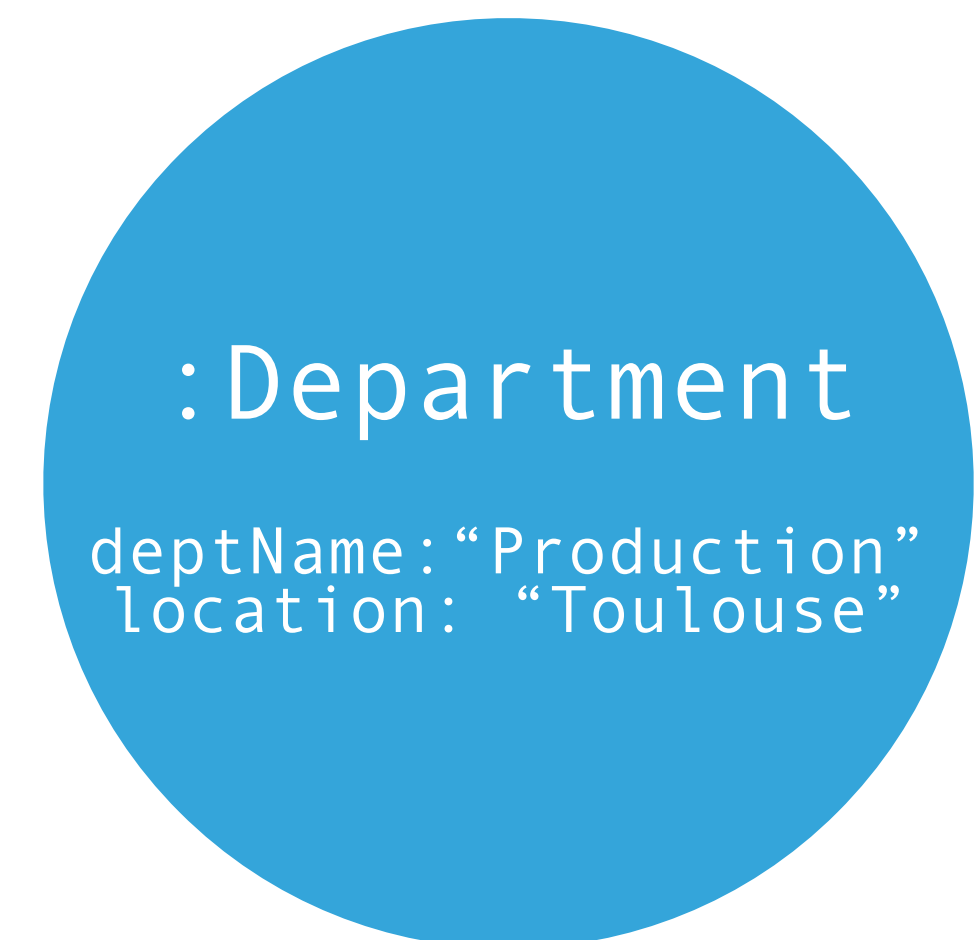
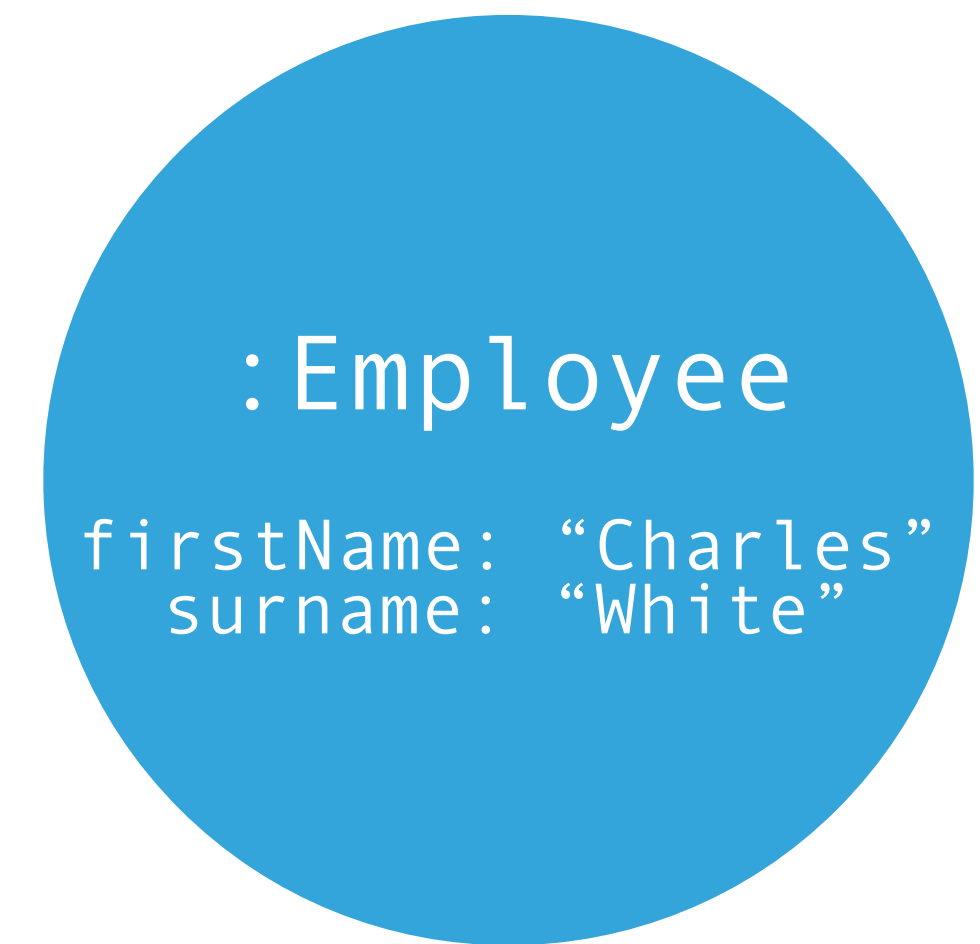
NEO4J DATA MODEL: LABELED PROPERTY GRAPH

- ▶ **Labeled Property Graph**
- ▶ **Nodes** have *labels* (types)
- ▶ **Nodes** can have *properties* (name-value pairs)
- ▶ **Edges** have a start node, an end node, and a direction
- ▶ **Edges** can also have *name* and *properties*



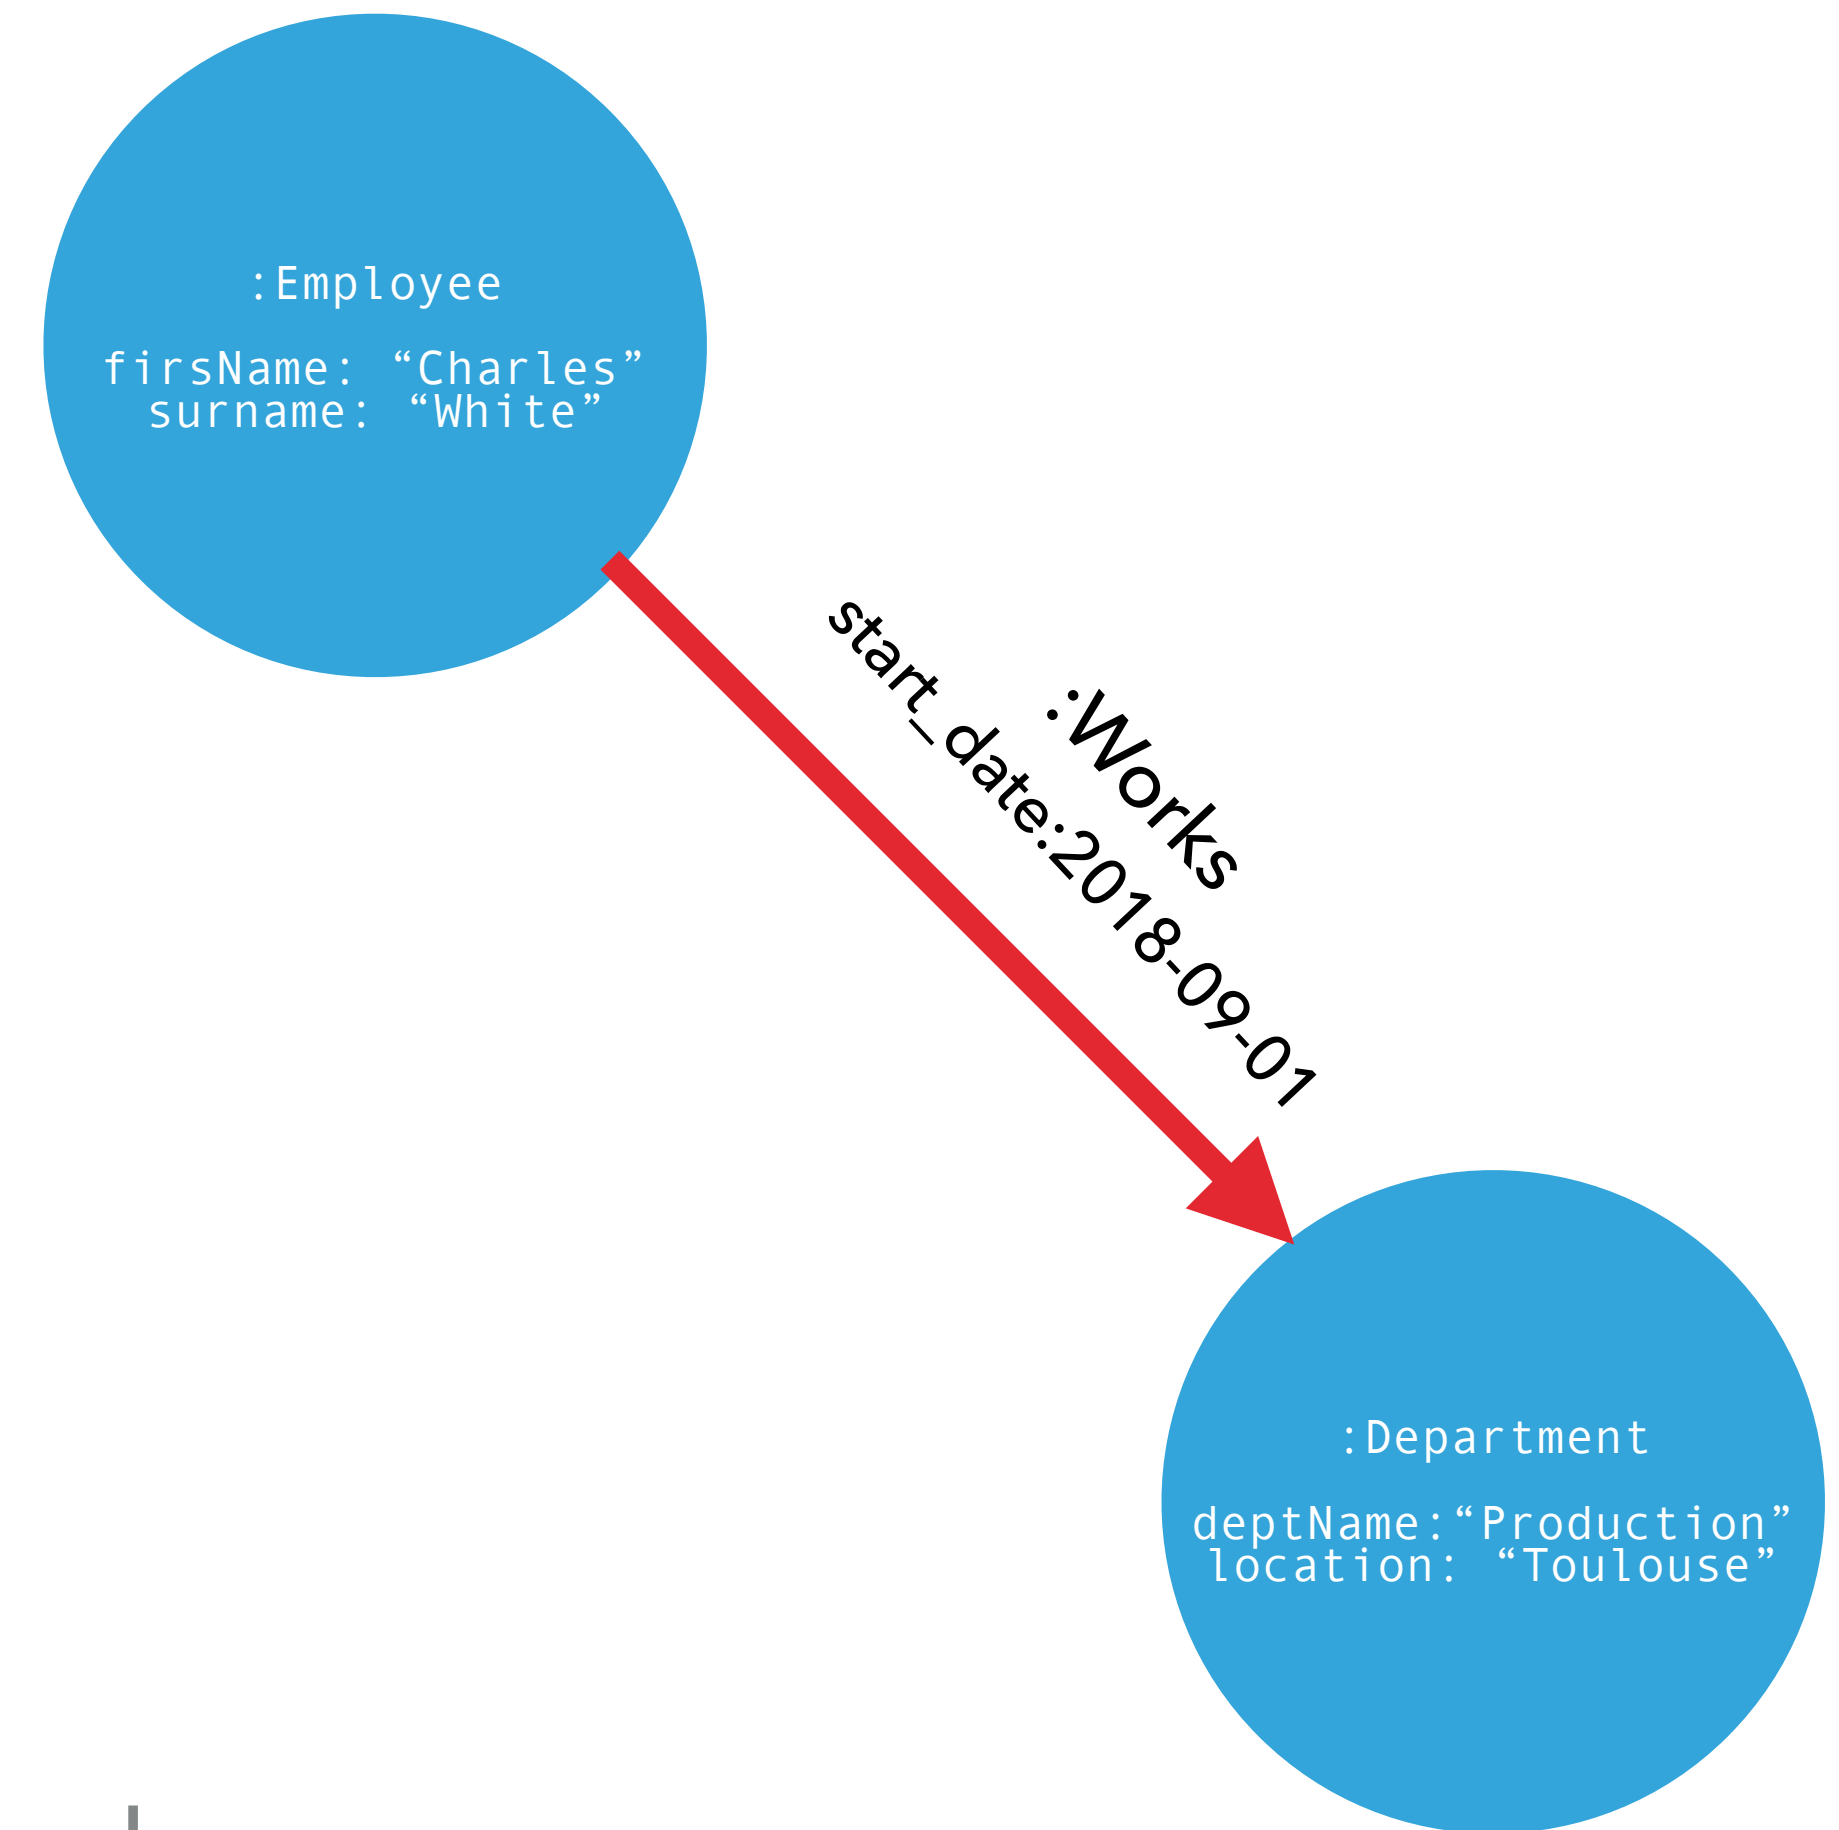
NODES IN NEO4J

- ▶ Represent **entities** with a unique conceptual identity
 - ▶ Have unique (internal) identifier
 - ▶ Can have **types**, defined by one or more **labels**
- ▶ **Labels**
 - ▶ Nodes having the same label belong to the same set
 - ▶ Can be added and removed at runtime
 - ▶ Naming convention: Camel case, beginning with an upper-case character
- ▶ **Properties**
 - ▶ **name:value** pairs (no lists or array – deprecated)
 - ▶ **name** is a string
 - ▶ Standard data types (Integer, Float, String, Boolean, spatial Point, DateTime)
 - ▶ Lower camel case, beginning with a lower-case character



RELATIONSHIPS

- ▶ It connects *source node* and an *target node*
 - ▶ Directions are **not** binding at query time
 - ▶ Loops are allowed
- ▶ Relationship **Type**
 - ▶ The name of the relationship
 - ▶ Upper case, using underscore to separate words
- ▶ **Properties**
 - ▶ name:value pairs - as for nodes
- ▶ Built in "*referential integrity*"
 - ▶ An existing relationship never points to a non-existing node
 - ▶ You cannot delete a node without deleting its associated relationships



EXAMPLE DATABASE: IMDB

- ▶ Node labels
 - ▶ One for each IMDB table
 - ▶ Properties as in the matching IMDB tables
- ▶ Relationship types
 - ▶ One for each referential integrity constraint in the relational version

labels(n)	type(r)
["CastInfo"]	"HAS_CAST"
["KindType"]	"IS_OF_KIND"
["AKATitle"]	"ALSO_KNOWN_AS"
["MovieKeyword"]	"HAS_KEYWORD"
["Title"]	"HAS_CAST"
["RoleType"]	"HAS_ROLE"
["CharName"]	"PERFORMED_AS"
["Person"]	"PARTICIPATES_IN"
["CastInfo"]	"PARTICIPATES_IN"
["AKAName"]	"ALSO_KNOWN_AS"
["CastInfo"]	"PERFORMED_AS"
["CastInfo"]	"HAS_ROLE"
["Title"]	"HAS_KEYWORD"
["Keyword"]	"IS"
["Title"]	"IS_OF_KIND"
["MovieKeyword"]	"IS"
["Person"]	"ALSO_KNOWN_AS"
["Title"]	"ALSO_KNOWN_AS"



QUERYING NEO4J

CYPHER

- ▶ **Declarative** graph query language
- ▶ Allows for expressive and efficient querying and updates
 - ▶ Inspired by SQL (query clauses) and SPARQL (pattern matching)
- ▶ Many features stem from improving on SQL pain points
 - ▶ e.g. JOINS, GROUP BY

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)

RETURN distinct p.name
```

Query: find the name of all the actors that played a part in the movie "The Matrix" produced in 1999

QUERY CLAUSES

- ▶ MATCH
 - ▶ Specifies **graph patterns** to be searched for
- ▶ WHERE
 - ▶ **Filtering** constraints on nodes and relationships
- ▶ RETURN
 - ▶ Defines **returned** data
- ▶ ORDER BY
 - ▶ Describes how results should be **ordered**
- ▶ UNION
 - ▶ Merges results from two or more queries
- ▶ WITH
 - ▶ Chains subsequent query parts and forwards results from one to the next

CYPHER PRINCIPLES


- ▶ Declarative
- ▶ Specification by example
- ▶ Patterns specified with ASCII-Art inspired syntax
 - ▶ Circles () for nodes
 - ▶ Arrows <—, --, —> and [] for relationships

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    )-[r:HAS_CAST]->(c:CastInfo)<-[s:PARTICIPATES_IN]-(p:Person)

RETURN distinct p.name
```

MATCH CLAUSE

- ▶ Search for sub-graphs of the data graph that match **all** the **patterns** in the query
 - ▶ Patterns can occur many times in the graph
- ▶ The query result are set of **variable bindings**
 - ▶ each variable has to be bound to a node/relationship
- ▶ Queries can return:
 - ▶ a **graph** (node and relationship instances)
 - ▶ a **table** (node and relationship properties)
- ▶ Results are unordered



labels(n)	type(r)
["CastInfo"]	"HAS_CAST"
["KindType"]	"IS_OF_KIND"
["AKATitle"]	"ALSO_KNOWN_AS"
["MovieKeyword"]	"HAS_KEYWORD"
["Title"]	"HAS_CAST"
["RoleType"]	"HAS_ROLE"
["CharName"]	"PERFORMED_AS"
["Person"]	"PARTICIPATES_IN"
["CastInfo"]	"PARTICIPATES_IN"
["AKAName"]	"ALSO_KNOWN_AS"
["CastInfo"]	"PERFORMED_AS"
["CastInfo"]	"HAS_ROLE"
["Title"]	"HAS_KEYWORD"
["Keyword"]	"IS"
["Title"]	"IS_OF_KIND"
["MovieKeyword"]	"IS"
["Person"]	"ALSO_KNOWN_AS"
["Title"]	"ALSO_KNOWN_AS"



NODE PATTERN

```
(variable:Label:...:Label{key: expression, ..., key: expression})
```

- ▶ Matches nodes
- ▶ **Variable:** used to refer/access nodes
- ▶ **Labels:** all node labels to be matched
- ▶ **Properties:** properties of the nodes to be matched
 - ▶ Order is not important

EXAMPLES

```
(t)
```

- ▶ Any node, referred by the variable `t`

```
(t:Title)
```

- ▶ Nodes with label `T i t l e`

```
(t:Title:Person)
```

- ▶ Nodes with labels `T i t l e` **and** `P e r s o n`

```
(t:Title{title:'The Matrix',production_year: 1999})
```

- ▶ Nodes with label `T i t l e` having properties with the specified values

PATH PATTERNS

```
()<-|- [variable:type|...|type variableLenght {key: expression, ..., key: expression}] -|->()
```

- ▶ Describe multiple nodes and one or more relationships among them
- ▶ A series of connected nodes is called a **path**
 - ▶ A traversal of part of the graph
 - ▶ Describes a single path, not a general sub-graph
 - ▶ Used as part of a query to specify patterns
- ▶ **Variable**: used to refer/access relationships
- ▶ **Type**: all relationship types to be matched
- ▶ **VariableLenght**: describe paths of arbitrary lengths (not just one relationship)
- ▶ **Properties**: properties of the relationships to be matched

EXAMPLES PATH PATTERNS /1

```
(t) - [r:HAS_CAST] - (c)
```

- ▶ Matches relationships of type `HAS_CAST` between all nodes

```
(t) -> (c)
```

- ▶ Matches all relationships from `t` (any node) to `c` (any node)

```
(t) -- (c)
```

- ▶ Matches relationships in any direction between `t` to `c`

```
(t:Title) -> (c)
```

- ▶ All relationships from nodes with label `Title`

EXAMPLES PATH PATTERNS /2

```
(t) <- [HAS_CAST] - (c)
```

```
(t) - [HAS_CAST | HAS_KEYWORD] -> (c)
```

```
(t) - [r] - (c)
```

```
(t) - [*1..5] - (c)
```

- ▶ Relationships of type HAS_CAST from c to t
- ▶ Matches relationships of type HAS_CAST or HAS_KEYWORD from t to c
- ▶ Binds the relationships to the variable r
- ▶ Variable length path, from 1 to 5 relationships, between t and c

EXAMPLES PATH PATTERNS /3

```
(t) - [*] -> (c)
```

```
(t) - [HAS_CAST] - (c {movie_id: '2543774'})
```

```
shortestPath((p1:Person) - [*..6] - ((p2:Person)))
```

```
size((t) --> () --> ())
```

- ▶ Variable path of any number of 5 relationships, from t to c
- ▶ A relationship of type HAS_CAST from a node t to a node c having declared property
- ▶ Find a single shortest path between any two nodes Person
- ▶ Count the paths matching the pattern

RETURN CLAUSE

- ▶ Defines what to include in the query result set
- ▶ Projection of node and relationship variables
 - ▶ Properties accessed via dot notation
 - ▶ Missing properties are treated as `null`
 - ▶ `*` returns all the variables
 - ▶ `AS` allows to explicitly (re)name results
- ▶ Literals, predicates, properties, functions can also be returned
- ▶ `DISTINCT`
 - ▶ retrieves only unique rows depending on the columns that have been selected to output

EXAMPLES RETURN /1

Returns distinct people name

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)
RETURN distinct p.name
```

Returns distinct people name, a true/false/null value if they are female, and a constant literal

```
RETURN distinct p.name, p.gender = 'f', 'A Literal'
```

Renaming and expressions

```
RETURN t.production_year-2000 AS Millenium_Age
```

p.name
"Paterson, Owen"
"Yuen, Woo-Ping"
"Yuen, Shun-Yee"
"Zhang, Daxing"
"Worthy, Megan"
"Wrencher, Luke"
"Wrencher, Charly"
"Woodward, Lawrence"
"Whalley, Sinclair"
"Whittle, Chris"
"Walker, Richard"
"Vollmer, Justine"
"van Gyen, Marijke Rikki"
"Valcarce, Marcel"
"Varnes, Kevin"
"Tuella, Michelle"

EXAMPLES RETURN /2

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)
RETURN distinct t.title, p.name
```


Returns distinct movie name and people name

t.title	p.name
"The Matrix"	"Paterson, Owen"
"The Matrix"	"Yuen, Woo-Ping"
"The Matrix"	"Yuen, Shun-Yee"
"The Matrix"	"Zhang, Daxing"
"The Matrix"	"Worthy, Megan"
"The Matrix"	"Wrencher, Luke"
"The Matrix"	"Wrencher, Charly"
"The Matrix"	"Woodward, Lawrence"
"The Matrix"	"Whalley, Sinclair"
"The Matrix"	"Whittle, Chris"
"The Matrix"	"Walker, Richard"
"The Matrix"	"Vollmer, Justine"
"The Matrix"	"van Gyen, Marijke Rikki"
"The Matrix"	"Valcarce, Marcel"
"The Matrix"	"Varnes, Kevin"
"The Matrix"	"Tuella, Michelle"

EXAMPLES RETURN /3

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)
RETURN t,p
```


Returns a collection of nodes



EXAMPLES RETURN /4

Returns a graph

```
MATCH (t:Title{
    title: 'The Matrix',
    production_year: 1999}
    ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)
RETURN t,r,c,s,p
```




EXAMPLES RETURN /4

```
MATCH (t:Title{
      title: 'The Matrix',
      production_year: 1999}
      ) - [r:HAS_CAST] -> (c:CastInfo) <- [s:PARTICIPATES_IN] - (p:Person)
RETURN t,r,c,s,p
```

t	r	c	s	p
<div><pre>{ "title_id": "4477062", "phonetic_code": "M3623", "kind_id": "1", "title": "The Matrix", "production_year": 1999 }</pre></div>	<div><pre>{ }</pre></div>	<div><pre>{ "cast_id": "63195892", "movie_id": "4477062", "role_id": "11", "person_id": "1833362" }</pre></div>	<div><pre>{ }</pre></div>	<div><pre>{ "name": "Paterson, Owen", "name_pcode_cf": "P3625", "gender": "m", "name_pcode_nf": "05136", "imdb_index": "I", "person_id": "1833362" }</pre></div>
<div><pre>{ "title_id": "4477062", "phonetic_code": "M3623", "kind_id": "1", "title": "The Matrix", "production_year": 1999 }</pre></div>	<div><pre>{ }</pre></div>	<div><pre>{ "note": "(kung fu choreographer) (as Yuen Wo Ping)", "cast_id": "62758035", "movie_id": "4477062", "role_id": "10", }</pre></div>	<div><pre>{ }</pre></div>	<div><pre>{ "name": "Yuen, Woo-Ping", "surname_pcode": "Y5", "name_pcode_cf": "Y5152", "gender": "m", "name_pcode_nf": "W1525", }</pre></div>

EXAMPLES RETURN /5

```
MATCH (t:Title{
    title: 'The Matrix'}
    )-[r:HAS_CAST]->(c:CastInfo)<-[s:PARTICIPATES_IN]-(p:Person)
RETURN t,r,c,s,p
```



Returns multiple sub-graphs
(potentially disconnected)

WHERE CLAUSE /1


- ▶ Add constraints to the pattern to be matched
 - ▶ **Evaluation is at matching time**
- ▶ Usual operations
 - ▶ Checking on existing properties
 - ▶ Range (<, <=, >, >=)
 - ▶ String matching (starts with, ends with, contains, regular expressions)
- ▶ Boolean Operations (AND, OR, XOR, NOT)
- ▶ Filters on node labels, node properties, relationship types, relationship properties

EXAMPLES WHERE / CHECKING ON PROPERTIES, BOOLEAN OPERATIONS

```
MATCH (p:Person)
WHERE p.name_pcode_cf='A1652' AND p.name_pcode_nf='J2165'
RETURN p.name, p.gender
```

```
MATCH (p:Person{
    name_pcode_cf: 'A1652',
    name_pcode_nf: 'J2165'})
RETURN p.name, p.gender
```

```
MATCH (p:Person)
WHERE p.name_pcode_cf='A1652'
      AND (p.gender='m' OR p.gender='f')
RETURN p.name, p.gender
```



► Same semantics
(same results set)

► No results with null

EXAMPLES WHERE / STRING MATCHING WITH REGEX

```
MATCH (t:Title)
WHERE t.title=~ 'S.*'
RETURN t.title, t.production_year
```

- ▶ Titles having names starting with **S**

```
MATCH (t:Title)
WHERE t.title=~ '.r.*n'
RETURN t.title, t.production_year
```

- ▶ Titles with **r** as second letter and ending with **n**

```
MATCH (t:Title)
WHERE t.title=~ '.*Matrix.*'
RETURN t.title, t.production_year
```

- ▶ Titles containing the word **Matrix**

```
MATCH (t:Title)
WHERE t.name_pcode_cf=~ '.2*'
RETURN t.title, t.production_year
```

- ▶ Two characters, second is a 2

WHERE CLAUSE /2

- ▶ Path patterns can also be predicates!
 - ▶ true if at least one solution is found
 - ▶ No new variables
- ▶ Path patterns in MATCH behave differently than path patterns in WHERE
 - ▶ MATCH: patterns produce subgraph for every found path
 - ▶ WHERE: it filters any matched sub-path where connected nodes have no relationship

```
MATCH (t:Title)
WHERE t.title=~ '.*Matrix.*' AND (t)-[:HAS_CAST]-({note:'(creator)'})
RETURN t.title, t.production_year
```

MISSING VALUES AND NULL

- ▶ Missing Values
 - ▶ A missing property evaluates to `null`
 - ▶ Checking for it in a node will make the condition evaluate to `false`

- ▶ `null` has the same semantics and properties as in SQL
 - ▶ `null` is not equal to `null`
 - ▶ `IS NULL / IS NOT NULL` work like in SQL

EXAMPLES NULL

```
MATCH (p:Person)
WHERE p.imdb_index IS NULL
RETURN p.name, p.gender
```

```
MATCH (t:Title)
WHERE t.production_year > 2000 OR t.production_year IS NULL
RETURN t.title, t.production_year
```


AGGREGATION FUNCTIONS

- ▶ Aggregation functions are analogous to SQL's GROUP BY
 - ▶ Compute over all matching subgraphs
 - ▶ Non aggregate-expressions are used as group key
 - ▶ Matching sub-graphs are divided in buckets, according to the key value

```
RETURN t, count(*)
```

- ▶ Classic functions (avg, count, max, min, sum)
- ▶ collect returns a list containing the values returned by an expression

EXAMPLE AGGREGATION FUNCTIONS

```
MATCH (cn:CharName{char_name:'Neo'}) - [PERFORMED_AS] - (c:CastInfo)
RETURN COUNT(cn)
```

COUNT(cn)

3

```
MATCH (t:Title)
WHERE t.title=~ '.*Matrix.*'
RETURN collect(t.title)
```

collect(t.title)

["The Matrix", "The Matrix", "The Matrix in Real Life", "The Matrix", "The Matrix", "Seal Matrix", "The Matrix", "The Matrix", "The Matrix Online", "The Matrix Reloaded: Car Chase", "The Matrix Reloaded: Teahouse Fight", "The Matrix: Follow the White Rabbit"]

EXAMPLE AGGREGATION FUNCTIONS / GROUP BY

```
MATCH (t:Title)-[r:HAS_CAST]->(c:CastInfo)<-[s:PARTICIPATES_IN]-(p:Person)
RETURN t.title, t.production_year, COUNT(c) AS Total
```

t.title	t.production_year	Total
"(#1.46)"	1998	26
"(#1.54)"	1998	26
"(#1.6)"	1998	50
"Princess Lucaj"	2012	17
"Into the Hot Zone"	2014	10
"Cut in the Gut"	2014	37
"Gun N Hide"	2012	28
"Hostage Standoff"	2012	29
"Escape from Bear Island"	2012	4
"Grizzly 911"	2012	3
"Lo mejor de Alaska y Mario"	2012	7
"El Chiringuito de Jugones"	2014	5
"El Clon"	2010	1
"El Club de Archi"	2012	7
"El club de la comedia"	2007	2
"El club de Los Tigritos"	1994	76
"El club"	2004	15
"El color de la pasión"	2014	2

THE WITH CLAUSE

- ▶ Behaves analogously to the RETURN clause
- ▶ Does not output anything to the user, just forwards the current result to the subsequent clause
 - ▶ The output of one part is passed as input to another
- ▶ Useful to filter on aggregate values (there is no HAVING clause in Cypher)
- ▶ But it can also be used to decompose queries into "sub-queries"

EXAMPLE WITH CLAUSE / HAVING

```
MATCH (t:Title)-[r:HAS_CAST]->(c:CastInfo)<-[s:PARTICIPATES_IN]-(p:Person)
WITH p.name AS actor, avg(t.production_year) AS average_acting_year
WHERE average_acting_year > 2000
RETURN actor, average_acting_year
```

actor	average_acting_year
"Rolón, Fernando"	1998.0
"Halac, Martín"	1997.59090909090909
"Cotta, Gustavo"	1999.0769230769229
"Wyszogrod, Iván"	1999.4285714285716
"Carus, Juan"	1998.0
"Sdrech, Enrique"	1992.16666666666665
"Cerretani, Marcia"	1992.16666666666665
"Álvarez, Daniel"	1998.5833333333333
"Dori, Yair"	1996.95833333333335
"Culell, Pablo"	2004.108108108108
"Calleja, Silvina"	1998.0

actor	average_acting_year
"Culell, Pablo"	2004.108108108108
"Guevara, Nacha"	2002.8
"Seefeld, Martín"	2009.7288135593224
"Ranni, Rodolfo"	2002.4864864864865
"Simpson, Jen"	2012.0
"Habeger, Brice"	2012.0
"Johnston, D.K."	2011.4
"Carew, Christy"	2012.0
"Pillifant, Tom"	2011.0
"Collier II, Michael"	2011.0
"Johnston, Katherine R."	2012.0

EXAMPLE WITH CLAUSE / “SUB QUERIES”



```
MATCH (t:Title)-[i:IS_OF_KIND]->(k:KindType{kind_name:'movie'})
RETURN t.title
```

```
MATCH (k:KindType{kind_name:'movie'})
WITH k.kind_id AS movie_kind
MATCH (t:Title)
WHERE t.kind_id = movie_kind
RETURN t.title
```

movie_kind
"movie"
"tv series"
"tv movie"
"video movie"
"tv mini series"
"video game"
"episode"
"short"
"tv short"

t.title
"#1137"
"#Blondepeopleproblems"
"#DoTheMileyCyrus"
"#ImHere - THE CALL"
"#MurderSelfie"
"#NonaHendryx Transformation"
"#Persianize"
"#SavingScott"
"#Selfie"
"#Ya"
"\$30"
"\$pent"
"\$tiffed or How I Learned to Deal with Dissapointment"
"'Orrible"
"'Til Death Do Us Part"
"'Til Death Do Us Partner"

UNION

- ▶ Combines the result of multiple queries
- ▶ Number and names of columns must be identical in all queries
- ▶ UNION ALL keeps duplicates

```
MATCH (p:Person)
RETURN p.name
UNION
MATCH (t:Title)
RETURN t.title
```

OPTIONAL MATCH

- ▶ When no solution is found, one specific solution with all the variables bound to NULL is generated
 - ▶ Pretty much a LEFT join
 - ▶ Either the whole pattern is matched, or nothing is matched

```
MATCH (c:CastInfo)<-[s:PARTICIPATES_IN]-(p:Person)
WHERE p.name = 'Reeves, Keanu'
OPTIONAL MATCH (c)-[pa:PERFORMED_AS]->(cn:CharName)
RETURN p.name, cn.char_name
```

p.name	cn.char_name
"Reeves, Keanu"	null
"Reeves, Keanu"	"Klaatu"
"Reeves, Keanu"	"Himself"
"Reeves, Keanu"	"Eric"
"Reeves, Keanu"	"Himself - Winner"
"Reeves, Keanu"	"Himself"
"Reeves, Keanu"	"Himself - Guest"
"Reeves, Keanu"	"Himself"
"Reeves, Keanu"	"Neo"
"Reeves, Keanu"	"Kevin Lomax"
"Reeves, Keanu"	null
"Reeves, Keanu"	"Himself - Guest"
"Reeves, Keanu"	"Donaka Mark"

ORDER BY

- ▶ By default, order of results is not defined
- ▶ Multiple criteria can be specified
 - ▶ Default direction is ASC
- ▶ Sorting is allowed only on properties
 - ▶ Not on nodes and relationships

```
MATCH (t:Title)
RETURN t
ORDER BY t.production_year DESC, t.title ASC
```

WAYS TO IMPROVE PERFORMANCE

- ▶ Use parameters instead of literals when possible
 - ▶ Cypher can re-use the queries
- ▶ Set upper limit for variable length pattern
- ▶ Return only the data you need
 - ▶ Avoid returning whole nodes and relationships
- ▶ Use **PROFILE** / **EXPLAIN** to analyse the performance of your queries

- ▶ Know/learn the types of questions to answer, then create new relationships that speed up those questions
- ▶ Use indexes

OTHER THINGS YOU CAN DO

- ▶ Depth and Breadth-First Search
- ▶ Path Finding (e.g. Dijkstra, A*)
- ▶ Triadic Closure
- ▶ In general, network properties (e.g. centrality, betweenness, etc.)

WRAPPING UP

TODAY WE COVERED

- ▶ Neo4J: data model and query language

END OF LECTURE